

# Kubernetes CLI Guide - Sections 17-21 (Continued)

## 17. StatefulSets (Continued)

```
bash
```

```
#!/bin/bash
# save as statefulset-demo.sh (continued from previous section)
echo ".flink StatefulSet Demo: Scaling and Persistence"

# Deploy the MySQL StatefulSet
kubectl apply -f mysql-statefulset.yaml

echo "⏳ Waiting for StatefulSet to be ready..."
kubectl wait --for=condition=ready pod -l app=mysql --timeout=300s

echo "🔍 Initial StatefulSet status:"
kubectl get statefulset mysql
kubectl get pods -l app=mysql

echo "💾 Checking persistent volumes:"
kubectl get pvc -l app=mysql

# Test data persistence by connecting to each MySQL instance
echo "🔎 Testing individual pod identity and persistence:"
for i in {0..2}; do
    echo "--- mysql-$i ---"
    kubectl exec mysql-$i -- mysql -u root -prootpassword -e "
        CREATE DATABASE IF NOT EXISTS pod_$i;
        USE pod_$i;
        CREATE TABLE IF NOT EXISTS test_table (id INT, pod_name VARCHAR(50));
        INSERT INTO test_table VALUES ($i, 'mysql-$i');
        SELECT * FROM test_table;
    " 2>/dev/null || echo "Pod mysql-$i not ready yet"
done

# Demonstrate ordered scaling
echo "📈 Testing ordered scaling (scale up to 5):"
kubectl scale statefulset mysql --replicas=5
kubectl get pods -l app=mysql -w &
WATCH_PID=$!
sleep 30
kill $WATCH_PID

echo "📝 Testing ordered scaling (scale down to 2):"
kubectl scale statefulset mysql --replicas=2
kubectl get pods -l app=mysql

# Test persistence after pod deletion
```

```
echo "⌚ Testing persistence after pod deletion:"  
kubectl delete pod mysql-1  
kubectl wait --for=condition=ready pod mysql-1 --timeout=120s  
  
# Verify data survived pod recreation  
kubectl exec mysql-1 -- mysql -u root -prootpassword -e "  
  USE pod_1;  
  SELECT * FROM test_table;  
" 2>/dev/null  
  
echo "✅ StatefulSet demo complete!"  
echo "🧹 Cleanup: kubectl delete statefulset mysql && kubectl delete service mysql-headless && kubectl delete pvc -l
```

## Mini-Project: Distributed Database Cluster

Create a complete distributed database setup with leader election and data replication:

bash

```

#!/bin/bash
# save as distributed-db-cluster.sh
echo "Creating Distributed Database Cluster"

# Create ConfigMap for database configuration
cat << EOF | kubectl apply -f -
apiVersion: v1
kind: ConfigMap
metadata:
  name: postgres-config
data:
  primary_init.sh: |
    #!/bin/bash
    set -e
    if [ ! -f /var/lib/postgresql/data/PG_VERSION ]; then
      echo "Initializing primary database..."
      initdb -D /var/lib/postgresql/data
      echo "host all all 0.0.0.0/0 md5" >> /var/lib/postgresql/data/pg_hba.conf
      echo "listen_addresses = '*' >> /var/lib/postgresql/data/postgresql.conf
      echo "wal_level = replica" >> /var/lib/postgresql/data/postgresql.conf
      echo "max_wal_senders = 3" >> /var/lib/postgresql/data/postgresql.conf
      echo "max_replication_slots = 3" >> /var/lib/postgresql/data/postgresql.conf
    fi
  replica_init.sh: |
    #!/bin/bash
    set -e
    if [ ! -f /var/lib/postgresql/data/PG_VERSION ]; then
      echo "Initializing replica database..."
      pg_basebackup -h postgres-0.postgres-headless -D /var/lib/postgresql/data -U postgres -v -P -W
      echo "standby_mode = 'on'" >> /var/lib/postgresql/data/recovery.conf
      echo "primary_conninfo = 'host=postgres-0.postgres-headless port=5432 user=postgres'" >> /var/lib/postgresql/
    fi
EOF

# Create headless service for StatefulSet
cat << EOF | kubectl apply -f -
apiVersion: v1
kind: Service
metadata:
  name: postgres-headless
  labels:
    app: postgres
spec:

```

```
ports:
- port: 5432
  name: postgres
clusterIP: None
selector:
  app: postgres
---
# Create regular service for client connections (points to primary)
apiVersion: v1
kind: Service
metadata:
  name: postgres-primary
labels:
  app: postgres
spec:
  ports:
  - port: 5432
    name: postgres
selector:
  app: postgres
  role: primary
EOF
```

*# Create the PostgreSQL StatefulSet with leader election*

```
cat << EOF | kubectl apply -f -
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: postgres
spec:
  serviceName: postgres-headless
  replicas: 3
  selector:
    matchLabels:
      app: postgres
  template:
    metadata:
      labels:
        app: postgres
  spec:
    containers:
    - name: postgres
      image: postgres:13
      env:
```

```
- name: POSTGRES_PASSWORD
  value: "secretpassword"
- name: POSTGRES_USER
  value: "postgres"
- name: POSTGRES_DB
  value: "testdb"
- name: PGDATA
  value: /var/lib/postgresql/data/pgdata
- name: POD_NAME
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
ports:
- containerPort: 8080
  name: http
  protocol: TCP
env:
- name: ENVIRONMENT
  value: ${ENVIRONMENT}
- name: LOG_LEVEL
  value: ${LOG_LEVEL}
- name: DB_PASSWORD
  valueFrom:
    secretKeyRef:
      name: ${APP_NAME}-secrets
      key: database-password
- name: API_KEY
  valueFrom:
    secretKeyRef:
      name: ${APP_NAME}-secrets
      key: api-key
envFrom:
- configMapRef:
    name: ${APP_NAME}-config
resources:
requests:
  memory: ${MEMORY_REQUEST}
  cpu: ${CPU_REQUEST}
limits:
  memory: ${MEMORY_LIMIT}
  cpu: ${CPU_LIMIT}
livenessProbe:
  httpGet:
    path: /health
```

```
    port: 8080
    initialDelaySeconds: 30
    periodSeconds: 10
    timeoutSeconds: 5
    failureThreshold: 3
  readinessProbe:
    httpGet:
      path: /ready
      port: 8080
    initialDelaySeconds: 5
    periodSeconds: 5
    timeoutSeconds: 3
    failureThreshold: 3
  volumeMounts:
    - name: app-config
      mountPath: /etc/config
      readOnly: true
    - name: temp-storage
      mountPath: /tmp
  volumes:
    - name: app-config
      configMap:
        name: ${APP_NAME}-config
    - name: temp-storage
      emptyDir: {}
  affinity:
    podAntiAffinity:
      preferredDuringSchedulingIgnoredDuringExecution:
        - weight: 100
          podAffinityTerm:
            labelSelector:
              matchExpressions:
                - key: app
                  operator: In
                  values:
                    - ${APP_NAME}
    topologyKey: kubernetes.io/hostname
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: ${APP_NAME}-service
  namespace: ${APP_NAME}-${ENVIRONMENT}
  labels:
```

```
app: ${APP_NAME}
spec:
type: ClusterIP
ports:
- port: 80
  targetPort: 8080
  protocol: TCP
  name: http
selector:
app: ${APP_NAME}

---
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
name: ${APP_NAME}-ingress
namespace: ${APP_NAME}-${ENVIRONMENT}
annotations:
nginx.ingress.kubernetes.io/rewrite-target: /
nginx.ingress.kubernetes.io/ssl-redirect: "true"
cert-manager.io/cluster-issuer: "letsencrypt-prod"
spec:
tls:
- hosts:
- ${APP_NAME}-${ENVIRONMENT}.${DOMAIN}
secretName: ${APP_NAME}-tls
rules:
- host: ${APP_NAME}-${ENVIRONMENT}.${DOMAIN}
http:
paths:
- path: /
  pathType: Prefix
backend:
service:
name: ${APP_NAME}-service
port:
number: 80

---
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
name: ${APP_NAME}-hpa
namespace: ${APP_NAME}-${ENVIRONMENT}
spec:
scaleTargetRef:
```

```
apiVersion: apps/v1
kind: Deployment
  name: ${APP_NAME}
  minReplicas: ${MIN_REPLICAS}
  maxReplicas: ${MAX_REPLICAS}
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
    - type: Resource
      resource:
        name: memory
        target:
          type: Utilization
          averageUtilization: 80
---
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: ${APP_NAME}-pdb
  namespace: ${APP_NAME}-${ENVIRONMENT}
spec:
  minAvailable: 50%
  selector:
    matchLabels:
      app: ${APP_NAME}
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: ${APP_NAME}-network-policy
  namespace: ${APP_NAME}-${ENVIRONMENT}
spec:
  podSelector:
    matchLabels:
      app: ${APP_NAME}
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
```

```
- namespaceSelector:  
  matchLabels:  
    name: ingress-nginx  
- namespaceSelector:  
  matchLabels:  
    environment: ${ENVIRONMENT}  
ports:  
- protocol: TCP  
  port: 8080  
egress:  
- to: []  
ports:  
- protocol: TCP  
  port: 53  
- protocol: UDP  
  port: 53  
- to:  
- namespaceSelector:  
  matchLabels:  
    name: database  
ports:  
- protocol: TCP  
  port: 5432
```

EOF

```
# Create deployment pipeline script  
cat << 'EOF' > deploy-production-app.sh  
#!/bin/bash  
set -e  
  
# Configuration  
APP_NAME=${1:-"webapp"}  
ENVIRONMENT=${2:-"staging"}  
APP_VERSION=${3:-"latest"}  
CONFIG_FILE=${4:-"deployment-config.env"}  
  
echo "🚀 Production Deployment Pipeline"  
echo "===== "  
echo "Application: $APP_NAME"  
echo "Environment: $ENVIRONMENT"  
echo "Version: $APP_VERSION"  
  
# Load environment-specific configuration  
if [ -f "$CONFIG_FILE" ]; then
```

```

source "$CONFIG_FILE"
else
  echo "⚠ Configuration file not found: $CONFIG_FILE"
  echo "Creating default configuration..."

case $ENVIRONMENT in
  "development")
    REPLICAS=1
    MIN_REPLICAS=1
    MAX_REPLICAS=3
    CPU_REQUEST="100m"
    CPU_LIMIT="500m"
    MEMORY_REQUEST="128Mi"
    MEMORY_LIMIT="512Mi"
    LOG_LEVEL="debug"
    ;;
  "staging")
    REPLICAS=2
    MIN_REPLICAS=2
    MAX_REPLICAS=5
    CPU_REQUEST="200m"
    CPU_LIMIT="1000m"
    MEMORY_REQUEST="256Mi"
    MEMORY_LIMIT="1Gi"
    LOG_LEVEL="info"
    ;;
  "production")
    REPLICAS=3
    MIN_REPLICAS=3
    MAX_REPLICAS=10
    CPU_REQUEST="500m"
    CPU_LIMIT="2000m"
    MEMORY_REQUEST="512Mi"
    MEMORY_LIMIT="2Gi"
    LOG_LEVEL="warn"
    ;;
esac

# Default values
IMAGE_REGISTRY=${IMAGE_REGISTRY:-"registry.company.com"}
DOMAIN=${DOMAIN:-"company.com"}
DB_HOST=${DB_HOST:-"postgres.database.svc.cluster.local"}
DB_PORT=${DB_PORT:-"5432"}
CACHE_ENABLED=${CACHE_ENABLED:-"true"}

```

```
# Generate base64 encoded secrets (in production, use proper secret management)
DB_PASSWORD_B64=$(echo "secretpassword" | base64 -w 0)
API_KEY_B64=$(echo "api-key-12345" | base64 -w 0)
fi

# Pre-deployment validations
echo "🔍 Running pre-deployment validations..."

# Check if kubectl is configured
if ! kubectl cluster-info >/dev/null 2>&1; then
    echo "✗ kubectl not configured or cluster not accessible"
    exit 1
fi

# Check if the image exists (mock check)
echo "🌐 Validating container image: ${IMAGE_REGISTRY}/${APP_NAME}:${APP_VERSION}"
# In production, you would check the registry
# docker manifest inspect ${IMAGE_REGISTRY}/${APP_NAME}:${APP_VERSION} >/dev/null || {
#     echo "✗ Image not found in registry"
#     exit 1
# }

# Security scan (mock)
echo "🔒 Running security scan..."
# In production, integrate with tools like Twistlock, Aqua, or Snyk
sleep 2
echo "✅ Security scan passed"

# Generate deployment manifests
echo "📝 Generating deployment manifests..."
TEMP_DIR=$(mktemp -d)
envsubst < production-app-template.yaml > "$TEMP_DIR/deployment.yaml"

# Validate generated YAML
echo "✅ Validating generated manifests..."
kubectl apply --dry-run=client -f "$TEMP_DIR/deployment.yaml" >/dev/null

# Deploy application
echo "🚀 Deploying application..."

# Create/update namespace first
kubectl apply -f - <<EOF
apiVersion: v1
```

```

kind: Namespace
metadata:
  name: ${APP_NAME}-${ENVIRONMENT}
labels:
  environment: ${ENVIRONMENT}
  managed-by: deployment-pipeline
EOF

# Apply all resources
kubectl apply -f "${TEMP_DIR}/deployment.yaml"

# Wait for deployment to complete
echo "⏳ Waiting for deployment to complete..."
kubectl rollout status deployment/${APP_NAME} -n ${APP_NAME}-${ENVIRONMENT} --timeout=600s

# Verify deployment
echo "🔍 Verifying deployment..."

# Check pod status
READY_PODS=$(kubectl get deployment ${APP_NAME} -n ${APP_NAME}-${ENVIRONMENT} -o jsonpath='{.status.read
DESIRED_PODS=$(kubectl get deployment ${APP_NAME} -n ${APP_NAME}-${ENVIRONMENT} -o jsonpath='{.spec.repl

if [ "$READY_PODS" = "$DESIRED_PODS" ]; then
  echo "✅ All pods are ready ($READY_PODS/$DESIRED_PODS)"
else
  echo "❌ Not all pods are ready ($READY_PODS/$DESIRED_PODS)"
  kubectl get pods -n ${APP_NAME}-${ENVIRONMENT}
  exit 1
fi

# Health check
echo "📋 Running health checks..."
kubectl run health-check-$(date +%s) \
  --image=busybox \
  --rm -i --restart=Never \
  --namespace=${APP_NAME}-${ENVIRONMENT} \
  -- wget -q --spider http://${APP_NAME}-service/health || {
  echo "❌ Health check failed"
  exit 1
}

echo "✅ Health check passed"

# Performance test (basic)

```

```
echo "⚡️ Running basic performance test..."  
kubectl run perf-test-$(date +%s) \  
  --image=busybox \  
  --rm -i --restart=Never \  
  --namespace=${APP_NAME}-$ENVIRONMENT \  
  -- sh -c "  
    for i in $(seq 1 10); do  
      wget -q -O http://${APP_NAME}-service/ >/dev/null && echo 'Request \$i: OK' || echo 'Request \$i: FAILED'  
    done  
  " || {  
    echo "⚠️ Performance test had some failures"  
}
```

# Generate deployment report

```
echo "📊 Generating deployment report..."  
REPORT_FILE="/tmp/deployment-report-${APP_NAME}-$ENVIRONMENT-$(date +%Y%m%d-%H%M%S).txt"
```

```
cat << REPORT > "$REPORT_FILE"
```

Deployment Report

=====

Application: \$APP\_NAME

Environment: \$ENVIRONMENT

Version: \$APP\_VERSION

Timestamp: \$(date)

Deployed by: \$(whoami)

Cluster: \$(kubectl config current-context)

Resources Created:

```
$(kubectl get all -n ${APP_NAME}-$ENVIRONMENT | grep -v "pod/health-check|pod/perf-test")
```

Pod Status:

```
$(kubectl get pods -n ${APP_NAME}-$ENVIRONMENT)
```

Service Status:

```
$(kubectl get service ${APP_NAME}-service -n ${APP_NAME}-$ENVIRONMENT)
```

Ingress Status:

```
$(kubectl get ingress ${APP_NAME}-ingress -n ${APP_NAME}-$ENVIRONMENT)
```

HPA Status:

```
$(kubectl get hpa ${APP_NAME}-hpa -n ${APP_NAME}-$ENVIRONMENT)
```

Recent Events:

```
$(kubectl get events -n ${APP_NAME}-$ENVIRONMENT --sort-by=.lastTimestamp | tail -10)
```

## REPORT

```
echo "📄 Deployment report saved: $REPORT_FILE"

# Cleanup
rm -rf "$TEMP_DIR"

echo ""
echo "🎉 Deployment completed successfully!"
echo ""
echo "📋 Next Steps:"
echo " Monitor: kubectl get pods -n ${APP_NAME}-${ENVIRONMENT} -w"
echo " Logs: kubectl logs -f deployment/${APP_NAME} -n ${APP_NAME}-${ENVIRONMENT}"
echo " Scale: kubectl scale deployment ${APP_NAME} --replicas=5 -n ${APP_NAME}-${ENVIRONMENT}"
echo " Rollback: kubectl rollout undo deployment/${APP_NAME} -n ${APP_NAME}-${ENVIRONMENT}"
echo ""
echo "🌐 Application URL: https://${APP_NAME}-${ENVIRONMENT}.${DOMAIN}"
EOF
```

```
chmod +x deploy-production-app.sh
```

```
# Create sample configuration files
cat << 'EOF' > deployment-config-dev.env
# Development Environment Configuration
ENVIRONMENT=development
REPLICAS=1
MIN_REPLICAS=1
MAX_REPLICAS=3
CPU_REQUEST="100m"
CPU_LIMIT="500m"
MEMORY_REQUEST="128Mi"
MEMORY_LIMIT="512Mi"
LOG_LEVEL="debug"
IMAGE_REGISTRY="registry.company.com"
DOMAIN="dev.company.com"
DB_HOST="postgres-dev.database.svc.cluster.local"
DB_PORT="5432"
CACHE_ENABLED="false"
DB_PASSWORD_B64="ZGV2cGFzc3dvcmQ=" # devpassword
API_KEY_B64="ZGV2LWFwaS1rZXk=" # dev-api-key
EOF
```

```
cat << 'EOF' > deployment-config-prod.env
# Production Environment Configuration
```

```
ENVIRONMENT=production
REPLICAS=3
MIN_REPLICAS=3
MAX_REPLICAS=10
CPU_REQUEST="500m"
CPU_LIMIT="2000m"
MEMORY_REQUEST="512Mi"
MEMORY_LIMIT="2Gi"
LOG_LEVEL="warn"
IMAGE_REGISTRY="registry.company.com"
DOMAIN="company.com"
DB_HOST="postgres-prod.database.svc.cluster.local"
DB_PORT="5432"
CACHE_ENABLED="true"
DB_PASSWORD_B64="cHJvZHBhc3N3b3Jk" # prodpassword
API_KEY_B64="cHJvZC1hcGkta2V5" # prod-api-key
EOF
```

```
echo "✅ Production deployment pipeline created!"
echo ""
echo "📁 Files created:"
echo " - production-app-template.yaml (Kubernetes manifests template)"
echo " - deploy-production-app.sh (Deployment script)"
echo " - deployment-config-dev.env (Development configuration)"
echo " - deployment-config-prod.env (Production configuration)"
echo ""
echo "📌 Usage Examples:"
echo " ./deploy-production-app.sh webapp development v1.0.0 deployment-config-dev.env"
echo " ./deploy-production-app.sh webapp production v1.0.0 deployment-config-prod.env"
```

## Mini-Project: Complete Kubernetes Operations Center

Create a comprehensive operations center with monitoring, alerting, and automated remediation:

```
bash
```

```

#!/bin/bash
# save as k8s-operations-center.sh
echo "kubelet Kubernetes Operations Center Setup"

# Create operations center directory structure
mkdir -p k8s-ops-center/{monitoring,alerting,automation,dashboards,reports}
cd k8s-ops-center

# Create comprehensive monitoring system
cat << 'EOF' > monitoring/cluster-monitor.sh
#!/bin/bash

METRICS_DIR="/tmp/k8s-metrics"
ALERT_THRESHOLD_CPU=80
ALERT_THRESHOLD_MEMORY=85
ALERT_THRESHOLD_PODS=90

mkdir -p "$METRICS_DIR"

collect_metrics() {
    local timestamp=$(date +%s)
    local date_str=$(date '+%Y-%m-%d %H:%M:%S')

    echo "📊 Collecting metrics at $date_str"

    # Node metrics
    if kubectl top nodes >/dev/null 2>&1; then
        kubectl top nodes --no-headers > "$METRICS_DIR/nodes-$timestamp.txt"
        kubectl top pods --all-namespaces --no-headers > "$METRICS_DIR/pods-$timestamp.txt"
    fi

    # Cluster state metrics
    kubectl get nodes -o json > "$METRICS_DIR/nodes-state-$timestamp.json"
    kubectl get pods --all-namespaces -o json > "$METRICS_DIR/pods-state-$timestamp.json"
    kubectl get events --all-namespaces -o json > "$METRICS_DIR/events-$timestamp.json"

    # Resource utilization
    cat << METRICS > "$METRICS_DIR/cluster-summary-$timestamp.txt"
    Timestamp: $date_str
    Nodes: $(kubectl get nodes --no-headers | wc -l)
    Total Pods: $(kubectl get pods --all-namespaces --no-headers | wc -l)
    Running Pods: $(kubectl get pods --all-namespaces --no-headers | grep Running | wc -l)
    Failed Pods: $(kubectl get pods --all-namespaces --no-headers | grep -E "(Error|Failed|CrashLoopBackOff)" | wc -l)
    METRICS
}

```

```

Services: $(kubectl get services --all-namespaces --no-headers | wc -l)
Deployments: $(kubectl get deployments --all-namespaces --no-headers | wc -l)
METRICS

# Cleanup old metrics (keep last 24 hours)
find "$METRICS_DIR" -name "*.txt" -o -name "*.json" | head -n -100 | xargs rm -f 2>/dev/null || true
}

analyze_trends() {
    echo "📈 Analyzing trends..."

    # CPU usage trend
    if ls "$METRICS_DIR"/nodes-*.txt >/dev/null 2>&1; then
        echo "💻 Node CPU Usage Trend (last 10 readings):"
        ls -t "$METRICS_DIR"/nodes-*.txt | head -10 | while read file; do
            timestamp=$(basename "$file" .txt | cut -d'-' -f2)
            date_str=$(date -d "@$timestamp" '+%H:%M:%S' 2>/dev/null || echo "$timestamp")
            avg_cpu=$(awk '{sum+=$3} END {if(NR>0) print sum/NR; else print 0}' "$file" | cut -d'%' -f1)
            echo " $date_str: ${avg_cpu}% average CPU"
        done
    fi
}

generate_alerts() {
    echo "⚠️ Checking alert conditions..."

    # Check for high resource usage
    if [ -f "$METRICS_DIR/nodes-$(ls \"$METRICS_DIR\" | grep \"nodes-\" | tail -1 | cut -d'-' -f2)"; then
        while read node cpu_usage cpu_percent memory_usage memory_percent; do
            cpu_num=$(echo $cpu_percent | tr -d '%')
            mem_num=$(echo $memory_percent | tr -d '%')

            if [ "$cpu_num" -gt "$ALERT_THRESHOLD_CPU" ]; then
                echo "🔥 ALERT: High CPU on $node: $cpu_percent"
                # In production, send to alerting system
            fi

            if [ "$mem_num" -gt "$ALERT_THRESHOLD_MEMORY" ]; then
                echo "MemoryWarning on $node: $memory_percent"
                # In production, send to alerting system
            fi
        done < "$METRICS_DIR/nodes-$(ls \"$METRICS_DIR\" | grep \"nodes-\" | tail -1 | cut -d'-' -f2)"
    fi
}

```

```

# Main monitoring loop
case "$1" in
    "collect")
        collect_metrics
        ;;
    "analyze")
        analyze_trends
        ;;
    "alert")
        generate_alerts
        ;;
    "daemon")
        echo "⌚ Starting monitoring daemon..."
        while true; do
            collect_metrics
            analyze_trends
            generate_alerts
            sleep 60
        done
        ;;
    *)
        echo "Usage: $0 {collect|analyze|alert|daemon}"
        ;;
esac
EOF

```

**chmod +x monitoring/cluster-monitor.sh**

```

# Create automated remediation system
cat << 'EOF' > automation/auto-remediation.sh
#!/bin/bash

```

LOG\_FILE="/tmp/auto-remediation.log"

```

log_action() {
    echo "$(date '+%Y-%m-%d %H:%M:%S') - $1" | tee -a "$LOG_FILE"
}

```

```

restart_failing_pods() {
    log_action "⌚ Checking for failing pods..."

```

# Find pods in error states

kubectl get pods --all-namespaces --no-headers | grep -E "(Error|CrashLoopBackOff|ImagePullBackOff)" | while read

```

log_action "⚠️ Found failing pod: $namespace/$pod"

# Get pod age
age=$(kubectl get pod "$pod" -n "$namespace" -o jsonpath='{.metadata.creationTimestamp}')
current_time=$(date -u +%s)
pod_time=$(date -d "$age" +%s 2>/dev/null || echo $current_time)
age_minutes=$(( (current_time - pod_time) / 60 ))

# Only restart if pod has been failing for more than 5 minutes
if [ "$age_minutes" -gt 5 ]; then
    log_action "♻️ Restarting pod $namespace/$pod (age: ${age_minutes}m)"
    kubectl delete pod "$pod" -n "$namespace"

    # Wait and verify
    sleep 30
    new_status=$(kubectl get pod -l app=$(kubectl get pod "$pod" -n "$namespace" -o jsonpath='{.metadata.labels.app}'))
    log_action "✅ Pod restart result: $new_status"
else
    log_action "⏳ Pod is new, waiting before restart (age: ${age_minutes}m)"
fi
done
}

scale_deployments() {
log_action "📊 Checking deployment scaling needs..."

# Check HPA status and manually scale if needed
kubectl get hpa --all-namespaces --no-headers | while read namespace hpa target min max replicas rest; do
    current_replicas=$(echo $replicas | cut -d',' -f1)
    max_replicas=$(echo $replicas | cut -d',' -f2)

    # If at max capacity, log warning
    if [ "$current_replicas" = "$max_replicas" ]; then
        log_action "⚠️ HPA $namespace/$hpa at maximum capacity ($current_replicas/$max_replicas)"
    fi
done
}

cleanup_completed_jobs() {
log_action "🧹 Cleaning up completed jobs..."

# Clean up completed jobs older than 1 hour
kubectl get jobs --all-namespaces -o json | jq -r '
.items[] |

```

```

select(.status.conditions[]?.type == "Complete" and .status.conditions[]?.status == "True") |
select((now - (.status.completionTime | fromdateiso8601)) > 3600) |
"\(.metadata.namespace) \(.metadata.name)"
' | while read namespace job; do
  log_action "🗑️ Deleting completed job: $namespace/$job"
  kubectl delete job "$job" -n "$namespace"
done
}

cleanup_failed_pods() {
log_action "⚡️ Cleaning up failed pods..."

# Clean up pods in Failed state older than 30 minutes
kubectl get pods --all-namespaces --field-selector=status.phase=Failed -o json | jq -r '
.items[] |
select((now - (.metadata.creationTimestamp | fromdateiso8601)) > 1800) |
"\(.metadata.namespace) \(.metadata.name)"
' | while read namespace pod; do
  log_action "🗑️ Deleting failed pod: $namespace/$pod"
  kubectl delete pod "$pod" -n "$namespace"
done
}

check_disk_pressure() {
log_action "💽 Checking for disk pressure..."

# Check nodes for disk pressure
kubectl get nodes -o json | jq -r '
.items[] |
select(.status.conditions[]? | select(.type == "DiskPressure" and .status == "True")) |
.metadata.name
' | while read node; do
  log_action "⚠️ Disk pressure detected on node: $node"

# Trigger cleanup on the node
kubectl get pods --all-namespaces --field-selector spec.nodeName="$node" -o json | jq -r '
.items[] |
select(.status.phase == "Succeeded" or .status.phase == "Failed") |
"\(.metadata.namespace) \(.metadata.name)"
' | head -5 | while read namespace pod; do
  log_action "⚡️ Cleaning up pod on pressured node: $namespace/$pod"
  kubectl delete pod "$pod" -n "$namespace"
done
done
}

```

```

}

# Main remediation functions
case "$1" in
    "pods")
        restart_failing_pods
        ;;
    "scale")
        scale_deployments
        ;;
    "jobs")
        cleanup_completed_jobs
        ;;
    "failed")
        cleanup_failed_pods
        ;;
    "disk")
        check_disk_pressure
        ;;
    "all")
        restart_failing_pods
        scale_deployments
        cleanup_completed_jobs
        cleanup_failed_pods
        check_disk_pressure
        ;;
    "daemon")
        log_action "🤖 Starting auto-remediation daemon..."
        while true; do
            restart_failing_pods
            cleanup_completed_jobs
            cleanup_failed_pods
            check_disk_pressure
            sleep 300 # Run every 5 minutes
        done
        ;;
    *)
        echo "🤖 Auto-Remediation System"
        echo "Usage: $0 {pods|scale|jobs|failed|disk|all|daemon}"
        ;;
esac
EOF

```

[chmod +x automation/auto-remediation.sh](#)

```
# Create dashboard generator
cat << 'EOF' > dashboards/generate-dashboard.sh
#!/bin/bash

DASHBOARD_DIR="/tmp/k8s-dashboard"
mkdir -p "$DASHBOARD_DIR"

generate_html_dashboard() {
    local dashboard_file="$DASHBOARD_DIR/k8s-dashboard-$(date +%Y%m%d-%H%M%S).html"

    cat << 'HTML' > "$dashboard_file"
    <!DOCTYPE html>
    <html>
    <head>
        <title>Kubernetes Operations Dashboard</title>
        <meta http-equiv="refresh" content="30">
        <style>
            body {
                font-family: 'Segoe UI', Arial, sans-serif;
                margin: 0;
                padding: 20px;
                background: #f5f5f5;
            }
            .header {
                background: #2196F3;
                color: white;
                padding: 20px;
                text-align: center;
                margin: -20px -20px 20px -20px;
            }
            .metrics-grid {
                display: grid;
                grid-template-columns: repeat(auto-fit, minmax(300px, 1fr));
                gap: 20px;
                margin-bottom: 20px;
            }
            .metric-card {
                background: white;
                padding: 20px;
                border-radius: 8px;
                box-shadow: 0 2px 4px rgba(0,0,0,0.1);
            }
            .metric-title {

```

```
    font-size: 18px;
    font-weight: bold;
    margin-bottom: 10px;
    color: #333;
}
.metric-value {
    font-size: 24px;
    font-weight: bold;
    color: #2196F3;
}
.status-good { color: #4CAF50; }
.status-warning { color: #FF9800; }
.status-critical { color: #F44336; }
.resource-table {
    width: 100%;
    border-collapse: collapse;
    margin-top: 10px;
}
.resource-table th, .resource-table td {
    padding: 8px;
    text-align: left;
    border-bottom: 1px solid #ddd;
}
.resource-table th {
    background: #f0f0f0;
    font-weight: bold;
}
pre {
    background: #f8f8f8;
    padding: 10px;
    border-radius: 4px;
    overflow-x: auto;
    font-size: 12px;
}
.section {
    background: white;
    margin: 20px 0;
    padding: 20px;
    border-radius: 8px;
    box-shadow: 0 2px 4px rgba(0,0,0,0.1);
}
</style>
</head>
<body>
```

```
<div class="header">
  <h1>  Kubernetes Operations Dashboard</h1>
  <p>Last Updated: TIMESTAMP_PLACEHOLDER</p>
  <p>Cluster: CLUSTER_PLACEHOLDER</p>
</div>

<div class="metrics-grid">
  <div class="metric-card">
    <div class="metric-title">  Nodes</div>
    <div class="metric-value">NODES_COUNT</div>
    <div>NODES_READY ready</div>
  </div>

  <div class="metric-card">
    <div class="metric-title">  Pods</div>
    <div class="metric-value">PODS_TOTAL</div>
    <div class="status-good">PODS_RUNNING running</div>
    <div class="status-critical">PODS_FAILED failed</div>
  </div>

  <div class="metric-card">
    <div class="metric-title">  Deployments</div>
    <div class="metric-value">DEPLOYMENTS_COUNT</div>
    <div>DEPLOYMENTS_READY ready</div>
  </div>

  <div class="metric-card">
    <div class="metric-title">  Services</div>
    <div class="metric-value">SERVICES_COUNT</div>
    <div>Across NAMESPACES_COUNT namespaces</div>
  </div>
</div>

<div class="section">
  <h2>  Resource Usage</h2>
  <table class="resource-table">
    <thead>
      <tr>
        <th>Node</th>
        <th>CPU Usage</th>
        <th>Memory Usage</th>
        <th>Status</th>
      </tr>
    </thead>

```

```

<tbody>
    NODE_USAGE_ROWS
</tbody>
</table>
</div>

<div class="section">
    <h2>⚠️ Recent Alerts</h2>
    <pre>RECENT_ALERTS</pre>
</div>

<div class="section">
    <h2>📋 Recent Events</h2>
    <pre>RECENT_EVENTS</pre>
</div>

<div class="section">
    <h2>🔍 Problematic Pods</h2>
    <pre>PROBLEMATIC PODS</pre>
</div>

<div class="section">
    <h2>💾 Storage Status</h2>
    <pre>STORAGE_STATUS</pre>
</div>
</body>
</html>
HTML

```

```

# Collect current data
local timestamp=$(date '+%Y-%m-%d %H:%M:%S')
local cluster=$(kubectl config current-context)

# Basic metrics
local nodes_count=$(kubectl get nodes --no-headers | wc -l)
local nodes_ready=$(kubectl get nodes --no-headers | grep -c Ready)
local pods_total=$(kubectl get pods --all-namespaces --no-headers | wc -l)
local pods_running=$(kubectl get pods --all-namespaces --no-headers | grep -c Running)
local pods_failed=$(kubectl get pods --all-namespaces --no-headers | grep -E "(Error|Failed|CrashLoopBackOff)")
local deployments_count=$(kubectl get deployments --all-namespaces --no-headers | wc -l)
local deployments_ready=$(kubectl get deployments --all-namespaces --no-headers | awk '$3==4 {count++} END {print count}')
local services_count=$(kubectl get services --all-namespaces --no-headers | wc -l)
local namespaces_count=$(kubectl get namespaces --no-headers | wc -l)

```

```

# Resource usage
local node_usage_rows=""
if kubectl top nodes >/dev/null 2>&1; then
    node_usage_rows=$(kubectl top nodes --no-headers | while read node cpu_usage cpu_percent memory_usage mem_percent; do
        local status_class="status-good"
        local cpu_num=$(echo $cpu_percent | tr -d '%')
        local mem_num=$(echo $memory_usage | tr -d '%')

        if [ "$cpu_num" -gt 80 ] || [ "$mem_num" -gt 80 ]; then
            status_class="status-critical"
        elif [ "$cpu_num" -gt 60 ] || [ "$mem_num" -gt 60 ]; then
            status_class="status-warning"
        fi

        echo "<tr><td>$node</td><td class=\"$status_class\">$cpu_percent</td><td class=\"$status_class\">$mem_percent</td></tr>"
    done)
else
    node_usage_rows=<tr><td colspan="4">Metrics server not available</td></tr>
fi

# Recent alerts (from log file)
local recent_alerts=$(tail -10 /tmp/auto-remediation.log 2>/dev/null | grep ALERT || echo "No recent alerts")

# Recent events
local recent_events=$(kubectl get events --all-namespaces --sort-by='.lastTimestamp' | tail -10)

# Problematic pods
local problematic_pods=$(kubectl get pods --all-namespaces | grep -E "(Error|Failed|CrashLoopBackOff|Pending)" || echo "No problematic pods found")

# Storage status
local storage_status=$(kubectl get pv,pvc --all-namespaces 2>/dev/null || echo "No storage resources found")

# Replace placeholders
sed -i "s/TIMESTAMP_PLACEHOLDER/$timestamp/g" "$dashboard_file"
sed -i "s/CLUSTER_PLACEHOLDER/$cluster/g" "$dashboard_file"
sed -i "s/NODES_COUNT/$nodes_count/g" "$dashboard_file"
sed -i "s/NODES_READY/$nodes_ready/g" "$dashboard_file"
sed -i "s/PODS_TOTAL/$pods_total/g" "$dashboard_file"
sed -i "s/PODS_RUNNING/$pods_running/g" "$dashboard_file"
sed -i "s/PODS_FAILED/$pods_failed/g" "$dashboard_file"
sed -i "s/DEPLOYMENTS_COUNT/$deployments_count/g" "$dashboard_file"
sed -i "s/DEPLOYMENTS_READY/$deployments_ready/g" "$dashboard_file"
sed -i "s/SERVICES_COUNT/$services_count/g" "$dashboard_file"
sed -i "s/NAMESPACES_COUNT/$namespaces_count/g" "$dashboard_file"

```

```
sed -i "s|NODE_USAGE_ROWS|$node_usage_rows|g" "$dashboard_file"

# Handle multi-line replacements
echo "$recent_alerts" > /tmp/alerts.txt
echo "$recent_events" > /tmp/events.txt
echo "$problematic_pods" > /tmp/pods.txt
echo "$storage_status" > /tmp/storage.txt

sed -i "/RECENT_ALERTS/r /tmp/alerts.txt" "$dashboard_file"
sed -i "s/RECENT_ALERTS//g" "$dashboard_file"

sed -i "/RECENT_EVENTS/r /tmp/events.txt" "$dashboard_file"
sed -i "s/RECENT_EVENTS//g" "$dashboard_file"

sed -i "/PROBLEMATIC PODS/r /tmp/pods.txt" "$dashboard_file"
sed -i "s/PROBLEMATIC PODS//g" "$dashboard_file"

sed -i "/STORAGE_STATUS/r /tmp/storage.txt" "$dashboard_file"
sed -i "s/STORAGE_STATUS//g" "$dashboard_file"

# Clean up temp files
rm -f /tmp/alerts.txt /tmp/events.txt /tmp/pods.txt /tmp/storage.txt

echo "📊 Dashboard generated: $dashboard_file"

# Create symlink to latest
ln -sf "$dashboard_file" "$DASHBOARD_DIR/latest.html"
echo "🔗 Latest dashboard: $DASHBOARD_DIR/latest.html"
}

# Main execution
case "$1" in
"generate")
    generate_html_dashboard
;;
"serve")
    generate_html_dashboard
    echo "🌐 Starting simple HTTP server on port 8080..."
    cd "$DASHBOARD_DIR"
    python3 -m http.server 8080 2>/dev/null || python -m SimpleHTTPServer 8080
;;
"auto")
    echo "⌚ Starting auto-refresh dashboard..."
    while true; do
```

```
generate_html_dashboard
sleep 30
done
;;
*)
echo "📊 Dashboard Generator"
echo "Usage: $0 {generate|serve|auto}"
echo ""
echo "Commands:"
echo " generate - Generate dashboard once"
echo " serve   - Generate and serve on HTTP port 8080"
echo " auto    - Auto-refresh every 30 seconds"
;;
esac
EOF
```

```
chmod +x dashboards/generate-dashboard.sh
```

```
# Create report generator
cat << 'EOF' > reports/generate-reports.sh
#!/bin/bash
```

```
REPORTS_DIR="/tmp/k8s-reports"
mkdir -p "$REPORTS_DIR"
```

```
generate_security_report() {
local report_file="$REPORTS_DIR/security-report-$(date +%Y%m%d).txt"

cat << REPORT > "$report_file"
Kubernetes Security Report
=====
Generated: $(date)
Cluster: $(kubectl config current-context)
```

```
1. Pod Security Analysis
```

```
-----
```

```
REPORT
```

```
echo "Pods running as root:" >> "$report_file"
kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.spec.securityContext.runAsUser == 0 or .spec.containers[].securityContext.runAsUser == 0) |
"\(.metadata.namespace)\(.metadata.name)"
' >> "$report_file" || echo "None found" >> "$report_file"
```

```

echo -e "\nPods with privileged containers:" >> "$report_file"
kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.spec.containers[].securityContext.privileged == true) |
"\(.metadata.namespace)\(.metadata.name)"
' >> "$report_file" || echo "None found" >> "$report_file"

echo -e "\n2. RBAC Analysis" >> "$report_file"
echo "-----" >> "$report_file"
echo "Service Accounts:" >> "$report_file"
kubectl get serviceaccounts --all-namespaces | wc -l >> "$report_file"

echo -e "\nCluster Roles:" >> "$report_file"
kubectl get clusterroles | wc -l >> "$report_file"

echo -e "\n3. Network Policies" >> "$report_file"
echo "-----" >> "$report_file"
kubectl get networkpolicies --all-namespaces >> "$report_file" 2>/dev/null || echo "No network policies found" >> "$report_file"

echo -e "\n4. Secrets Analysis" >> "$report_file"
echo "-----" >> "$report_file"
echo "Total secrets:" >> "$report_file"
kubectl get secrets --all-namespaces | wc -l >> "$report_file"

echo "  Security report generated: $report_file"
}

generate_performance_report() {
local report_file="$REPORTS_DIR/performance-report-$(date +%Y%m%d).txt"

cat << REPORT > "$report_file"
Kubernetes Performance Report
=====
Generated: $(date)
Cluster: $(kubectl config current-context)

1. Resource Utilization
-----
REPORT

if kubectl top nodes >/dev/null 2>&1; then
  echo "Node Resource Usage:" >> "$report_file"
  kubectl top nodes >> "$report_file"

```

```

echo -e "\nTop CPU Consuming Pods:" >> "$report_file"
kubectl top pods --all-namespaces --sort-by=cpu | head -10 >> "$report_file"

echo -e "\nTop Memory Consuming Pods:" >> "$report_file"
kubectl top pods --all-namespaces --sort-by=memory | head -10 >> "$report_file"
else
  echo "Metrics server not available" >> "$report_file"
fi

echo -e "\n2. Pod Restart Analysis" >> "$report_file"
echo "-----" >> "$report_file"
kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.status.containerStatuses[]?.restartCount > 0) |
"\(.metadata.namespace)\(\.metadata.name): \(.status.containerStatuses[0].restartCount) restarts"
' >> "$report_file"

echo -e "\n3. Storage Performance" >> "$report_file"
echo "-----" >> "$report_file"
kubectl get pv -o custom-columns=NAME:.metadata.name,CAPACITY:.spec.capacity.storage,STATUS:.status.phase >> "$report_file"

echo "  Performance report generated: $report_file"
}

```

```

generate_inventory_report() {
  local report_file="$REPORTS_DIR/inventory-report-$(date +%Y%m%d).txt"

  cat << REPORT > "$report_file"
  Kubernetes Inventory Report
  =====
  Generated: $(date)
  Cluster: $(kubectl config current-context)

```

## 1. Cluster Summary

---

```

Nodes: $(kubectl get nodes --no-headers | wc -l)
Namespaces: $(kubectl get namespaces --no-headers | wc -l)
Pods: $(kubectl get pods --all-namespaces --no-headers | wc -l)
Services: $(kubectl get services --all-namespaces --no-headers | wc -l)
Deployments: $(kubectl get deployments --all-namespaces --no-headers | wc -l)
ConfigMaps: $(kubectl get configmaps --all-namespaces --no-headers | wc -l)
Secrets: $(kubectl get secrets --all-namespaces --no-headers | wc -l)

```

## 2. Node Details

---

REPORT

```
kubectl get nodes -o wide >> "$report_file"

echo -e "\n3. Namespace Resource Distribution" >> "$report_file"
echo "-----" >> "$report_file"
kubectl get pods --all-namespaces --no-headers | awk '{ns[$1]++} END {for(n in ns) print n": "ns[n]" pods"}' | sort >>

echo -e "\n4. Image Analysis" >> "$report_file"
echo "-----" >> "$report_file"
kubectl get pods --all-namespaces -o json | jq -r '.items[].spec.containers[].image' | sort | uniq -c | sort -nr | head -20 >>

echo "  Inventory report generated: $report_file"
}
```

```
generate_compliance_report() {
    local report_file="$REPORTS_DIR/compliance-report-$(date +%Y%m%d).txt"

    cat << REPORT > "$report_file"
    Kubernetes Compliance Report
    =====
    Generated: $(date)
    Cluster: $(kubectl config current-context)
```

## 1. Resource Limits Compliance

---

REPORT

```
echo "Pods without resource limits:" >> "$report_file"
kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.spec.containers[0].resources.limits == null) |
"\(.metadata.namespace)/\(.metadata.name)"
' >> "$report_file"

echo -e "\n2. Security Context Compliance" >> "$report_file"
echo "-----" >> "$report_file"
echo "Pods without security context:" >> "$report_file"
kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.spec.securityContext == null and .spec.containers[].securityContext == null) |
"\(.metadata.namespace)/\(.metadata.name)"
```

```

' >> "$report_file"

echo -e "\n3. Label Compliance" >> "$report_file"
echo "-----" >> "$report_file"
echo "Resources without recommended labels:" >> "$report_file"
kubectl get deployments --all-namespaces -o json | jq -r '
.items[] |
select(.metadata.labels.app == null or .metadata.labels.version == null) |
"\(.metadata.namespace)\(.metadata.name)"
' >> "$report_file"

echo "  Compliance report generated: $report_file"
}

# Main execution
case "$1" in
"security")
    generate_security_report
;;
"performance")
    generate_performance_report
;;
"inventory")
    generate_inventory_report
;;
"compliance")
    generate_compliance_report
;;
"all")
    generate_security_report
    generate_performance_report
    generate_inventory_report
    generate_compliance_report
    echo "  All reports generated in: $REPORTS_DIR"
;;
*)
    echo "  Report Generator"
    echo "Usage: $0 {security|performance|inventory|compliance|all}"
    echo ""
    echo "Commands:"
    echo " security - Generate security analysis report"
    echo " performance - Generate performance analysis report"
    echo " inventory - Generate cluster inventory report"
    echo " compliance - Generate compliance check report"
esac

```

```
echo " all      - Generate all reports"
::
esac
EOF

chmod +x reports/generate-reports.sh

# Create master operations script
cat << 'EOF' > k8s-ops.sh
#!/bin/bash

OPS_CENTER_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"

show_menu() {
    echo " <img alt='Kubernetes logo' data-bbox='138 331 155 345' style='vertical-align: middle; height: 1em;"/> Kubernetes Operations Center"
    echo "====="
    echo ""
    echo "Monitoring:"
    echo " 1) Start monitoring daemon"
    echo " 2) Collect metrics once"
    echo " 3) Analyze trends"
    echo ""
    echo "Automation:"
    echo " 4) Start auto-remediation daemon"
    echo " 5) Run remediation once"
    echo " 6) Clean up resources"
    echo ""
    echo "Dashboard:"
    echo " 7) Generate dashboard"
    echo " 8) Serve dashboard (HTTP)"
    echo " 9) Auto-refresh dashboard"
    echo ""
    echo "Reports:"
    echo " 10) Generate security report"
    echo " 11) Generate performance report"
    echo " 12) Generate all reports"
    echo ""
    echo " 0) Exit"
    echo ""
}

execute_choice() {
    case $1 in
        1)
```

```
echo "⌚ Starting monitoring daemon..."  
"$OPS_CENTER_DIR/monitoring/cluster-monitor.sh" daemon  
:::  
2)  
echo "📊 Collecting metrics..."  
"$OPS_CENTER_DIR/monitoring/cluster-monitor.sh" collect  
:::  
3)  
echo "📈 Analyzing trends..."  
"$OPS_CENTER_DIR/monitoring/cluster-monitor.sh" analyze  
:::  
4)  
echo "🤖 Starting auto-remediation daemon..."  
"$OPS_CENTER_DIR/automation/auto-remediation.sh" daemon  
:::  
5)  
echo "⚡️ Running remediation..."  
"$OPS_CENTER_DIR/automation/auto-remediation.sh" all  
:::  
6)  
echo "🧹 Cleaning up resources..."  
"$OPS_CENTER_DIR/automation/auto-remediation.sh" jobs  
"$OPS_CENTER_DIR/automation/auto-remediation.sh" failed  
:::  
7)  
echo "📊 Generating dashboard..."  
"$OPS_CENTER_DIR/dashboards/generate-dashboard.sh" generate  
:::  
8)  
echo "🌐 Starting dashboard server..."  
"$OPS_CENTER_DIR/dashboards/generate-dashboard.sh" serve  
:::  
9)  
echo "⌚ Starting auto-refresh dashboard..."  
"$OPS_CENTER_DIR/dashboards/generate-dashboard.sh" auto  
:::  
10)  
echo "🔒 Generating security report..."  
"$OPS_CENTER_DIR/reports/generate-reports.sh" security  
:::  
11)  
echo "⚡️ Generating performance report..."  
"$OPS_CENTER_DIR/reports/generate-reports.sh" performance  
:::
```

```

12)
    echo "📊 Generating all reports..."
    "$OPS_CENTER_DIR/reports/generate-reports.sh" all
    ;;

0)
    echo "👋 Goodbye!"
    exit 0
    ;;

*)
    echo "✖ Invalid choice. Please try again."
    ;;
esac
}

# Interactive mode
if [ "$1" = "" ]; then
    while true; do
        clear
        show_menu
        read -p "Enter your choice [0-12]: " choice
        echo ""
        execute_choice "$choice"
        echo ""
        read -p "Press Enter to continue..."
    done
else
    # Direct command mode
    execute_choice "$1"
fi
EOF

```

**chmod +x k8s-ops.sh**

```

# Create README for the operations center
cat << 'EOF' > README.md
# 📄 Kubernetes Operations Center

```

A comprehensive operations center **for** monitoring, managing, and maintaining Kubernetes clusters.

**## Features**

**### 🔍 Monitoring**

- Real-time cluster metrics collection
- Resource usage trend analysis

- Automated alerting **for** threshold breaches
- Historical data retention

### **Automation**

- Auto-remediation of failing pods
- Cleanup of completed **jobs** and failed resources
- Disk pressure management
- Scaling recommendations

### **Dashboards**

- Real-time HTML dashboard
- Resource utilization visualization
- Alert status display
- Auto-refreshing capabilities

### **Reports**

- Security compliance analysis
- Performance assessment
- Resource inventory
- Compliance checking

## *Quick Start*

```
```bash
# Interactive mode
./k8s-ops.sh

# Direct commands
./k8s-ops.sh 1 # Start monitoring daemon
./k8s-ops.sh 4 # Start auto-remediation
./k8s-ops.sh 7 # Generate dashboard
./k8s-ops.sh 12 # Generate all reports
```

## Directory Structure

```
k8s-ops-center/
├── monitoring/      # Monitoring scripts
├── automation/     # Auto-remediation scripts
├── dashboards/      # Dashboard generation
├── reports/         # Report generators
├── k8s-ops.sh       # Main operations script
└── README.md        # This file
```

# Components

## Monitoring System

- **cluster-monitor.sh:** Collects metrics and analyzes trends
- Stores metrics in `/tmp/k8s-metrics/`
- Provides alerting based on configurable thresholds

## Auto-Remediation

- **auto-remediation.sh:** Automated problem resolution
- Restarts failing pods
- Cleans up resources
- Manages disk pressure

## Dashboard

- **generate-dashboard.sh:** Creates HTML dashboards
- Real-time cluster status
- Resource utilization graphs
- Serves via HTTP on port 8080

## Reports

- **generate-reports.sh:** Comprehensive reporting
- Security analysis
- Performance metrics
- Compliance checking
- Resource inventory

## Configuration

Edit the scripts to customize:

- Alert thresholds
- Cleanup intervals
- Report formats
- Dashboard refresh rates

## Dependencies

- kubectl configured and working
- jq for JSON processing
- Python for HTTP server
- Standard Unix tools (awk, sed, grep)

## Best Practices

1. Run monitoring and auto-remediation as background services
2. Review reports regularly for trends
3. Customize thresholds based on your environment
4. Integrate with external alerting systems
5. Regular backup of configurations and data

## Troubleshooting

- Check kubectl connectivity: `kubectl cluster-info`
- Verify metrics server: `kubectl top nodes`
- Check log files in `/tmp/` directory
- Ensure proper RBAC permissions

## Support

For issues and improvements, check the script comments and customize for your environment.

EOF

cd ..

```
echo "✅ Kubernetes Operations Center Setup Complete!"  
echo ""  
echo "📁 Directory Structure:"  
find k8s-ops-center -type f | head -15  
  
echo ""  
echo "🚀 Getting Started:"  
echo " cd k8s-ops-center"  
echo " ./k8s-ops.sh           # Interactive mode"
```

```
echo " ./k8s-ops.sh 7          # Generate dashboard"
echo " ./k8s-ops.sh 12          # Generate all reports"

echo ""
echo " 🛡️ Operations Center Features:"
echo "   ✓ Real-time monitoring"
echo "   ✓ Auto-remediation"
echo "   ✓ Interactive dashboards"
echo "   ✓ Comprehensive reporting"
echo "   ✓ Security analysis"
echo "   ✓ Performance monitoring"

echo ""
echo " 📖 See k8s-ops-center/README.md for detailed documentation"
```

## ## Conclusion

This comprehensive Kubernetes CLI guide has covered everything from basic concepts to advanced production operations. You now have:

### ### 🚀 \*\*Core Skills\*\*

- Essential kubectl commands and workflows
- Pod, Service, and Deployment management
- Storage and networking configuration
- Security with RBAC and Secrets

### ### 🚀 \*\*Advanced Capabilities\*\*

- Helm chart development and management
- Multi-environment deployment strategies
- Resource optimization and monitoring
- Automated troubleshooting and remediation

### ### 🛠 \*\*Production-Ready Tools\*\*

- Complete CI/CD pipeline templates
- Operations center with monitoring and alerting
- Security and compliance reporting
- Performance analysis and optimization

### ### 📄 \*\*Best Practices\*\*

- Declarative configuration management
- Resource quotas and limits
- Health checks and monitoring
- Backup and disaster recovery

The examples, demos, and mini-projects provide hands-on experience with real-world scenarios you'll encounter in production Kubernetes environments. Each section builds upon previous knowledge while introducing new concepts and tools.

Remember to:

- Always test in development environments first
- Implement proper RBAC and security measures
- Monitor resource usage and set appropriate limits
- Maintain regular backups and documentation
- Stay updated with Kubernetes releases and best practices

This guide serves as both a learning resource and a reference for ongoing Kubernetes operations. The scripts and tools provided can be customized and extended to meet your specific requirements and environments.5432

```
name: postgres
volumeMounts:
- name: postgres-data
  mountPath: /var/lib/postgresql/data
- name: config
  mountPath: /scripts
lifecycle:
postStart:
  exec:
    command:
    - /bin/bash
    - --c
    - |
      if [[ "${POD_NAME}" == "postgres-0" ]]; then
        echo "Setting up as primary"
        kubectl label pod ${POD_NAME} role=primary --overwrite
        /scripts/primary_init.sh
      else
        echo "Setting up as replica"
        kubectl label pod ${POD_NAME} role=replica --overwrite
        /scripts/replica_init.sh
      fi
resources:
requests:
  memory: 256Mi
  cpu: 250m
limits:
  memory: 512Mi
  cpu: 500m
volumes:
- name: config
configMap:
  name: postgres-config
  defaultMode: 0755
volumeClaimTemplates:
- metadata:
  name: postgres-data
spec:
  accessModes: [ "ReadWriteOnce" ]
resources:
  requests:
    storage: 5Gi
EOF
```

```

echo "⏳ Waiting for StatefulSet deployment..."
kubectl wait --for=condition=ready pod -l app=postgres --timeout=300s

echo "🔍 Checking cluster status:"
kubectl get pods -l app=postgres --show-labels
kubectl get pvc -l app=postgres

echo "📝 Testing database cluster:"
# Test primary database
echo "Testing primary (postgres-0):"
kubectl exec postgres-0 -- psql -U postgres -c "CREATE TABLE IF NOT EXISTS test_data (id SERIAL, data TEXT, created_at TIMESTAMP DEFAULT NOW());"
kubectl exec postgres-0 -- psql -U postgres -c "INSERT INTO test_data (data) VALUES ('Primary data entry');"

# Test replica access (after replication setup)
echo "Testing replica (postgres-1):"
kubectl exec postgres-1 -- psql -U postgres -c "SELECT * FROM test_data;" 2>/dev/null || echo "Replica not ready yet"

echo "✅ Distributed database cluster deployed!"
echo "📊 Monitor with:"
echo "  kubectl logs -f postgres-0"
echo "  kubectl exec postgres-0 -- psql -U postgres -c 'SELECT * FROM pg_stat_replication;'"
echo "✗ Cleanup:"
echo "  kubectl delete statefulset postgres"
echo "  kubectl delete service postgres-headless postgres-primary"
echo "  kubectl delete configmap postgres-config"
echo "  kubectl delete pvc -l app=postgres"

```

## 18. Resource Management

Resource management ensures efficient cluster utilization and prevents resource starvation. Understanding requests, limits, and quotas is crucial for production deployments.

### 18.1 Resource Requests and Limits

yaml

```
# resource-management-example.yaml
apiVersion: v1
kind: Pod
metadata:
  name: resource-demo
spec:
  containers:
    - name: app
      image: nginx
  resources:
    requests: # Minimum resources guaranteed to container
      memory: "64Mi" # 64 Megabytes of RAM
      cpu: "250m" # 250 millicores (0.25 CPU cores)
    limits: # Maximum resources container can use
      memory: "128Mi" # Container killed if it exceeds this
      cpu: "500m" # Container throttled if it exceeds this
    # Best Practice: Always set both requests and limits
    # Requests help scheduler make placement decisions
    # Limits prevent containers from consuming excessive resources
```

## Example: Quality of Service Classes

Understanding how Kubernetes assigns QoS classes based on resource specifications:

```
bash
```

```
#!/bin/bash
# save as qos-demo.sh
echo "🌐 Quality of Service Classes Demo"

# Create pods with different QoS classes

# Guaranteed QoS - requests = limits
cat << EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: guaranteed-pod
spec:
  containers:
    - name: app
      image: nginx
  resources:
    requests:
      memory: "256Mi"
      cpu: "500m"
    limits:
      memory: "256Mi"
      cpu: "500m"
EOF
```

```
# Burstable QoS - requests < limits or only requests specified
cat << EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: burstable-pod
spec:
  containers:
    - name: app
      image: nginx
  resources:
    requests:
      memory: "128Mi"
      cpu: "250m"
    limits:
      memory: "256Mi"
      cpu: "500m"
EOF
```

```

# BestEffort QoS - no requests or limits specified
cat << EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: besteffort-pod
spec:
  containers:
  - name: app
    image: nginx
EOF

echo "⌚️ Waiting for pods to be ready..."
kubectl wait --for=condition=ready pod --all --timeout=60s

echo "🔍 Checking QoS classes:"
for pod in guaranteed-pod burstable-pod besteffort-pod; do
  qos=$(kubectl get pod $pod -o jsonpath='{.status.qosClass}')
  echo "$pod: $qos"
done

echo "📊 Resource allocation:"
kubectl describe pods | grep -A 10 -B 2 "QoS Class"

echo "💡 QoS Class Implications:"
echo "Guaranteed: Highest priority, last to be evicted"
echo "Burstable: Medium priority, evicted before Guaranteed"
echo "BestEffort: Lowest priority, first to be evicted"

echo "🧹 Cleanup:"
echo "kubectl delete pod guaranteed-pod burstable-pod besteffort-pod"

```

## Demo: Resource Monitoring and Alerting

Create a comprehensive resource monitoring system:

bash

```

#!/bin/bash
# save as resource-monitor.sh
NAMESPACE=${1:-default}
THRESHOLD_CPU=80
THRESHOLD_MEMORY=80

echo "📊 Resource Monitoring for namespace: $NAMESPACE"

# Function to check if metrics-server is available
check_metrics_server() {
    if ! kubectl top nodes >/dev/null 2>&1; then
        echo "❌ Metrics server not available. Installing..."
        kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
        echo "⏳ Waiting for metrics server to start..."
        kubectl wait --for=condition=available deployment metrics-server -n kube-system --timeout=120s
    fi
}

# Function to analyze resource usage
analyze_resources() {
    echo "🔍 Current Resource Usage:"

    # Node-level analysis
    echo "--- Node Resources ---"
    kubectl top nodes

    # Pod-level analysis
    echo -e "\n--- Pod Resources in $NAMESPACE ---"
    kubectl top pods -n $NAMESPACE --sort-by=cpu

    # Check for resource-constrained pods
    echo -e "\n⚠️ Resource Alerts:"

    # Get pod resource usage and compare with limits
    kubectl get pods -n $NAMESPACE -o custom-columns=NAME:.metadata.name,CPU_REQUEST:.spec.containers[0].res

    # Check for pods without resource specifications
    echo -e "\n⚠️ Pods without resource limits:"
    kubectl get pods -n $NAMESPACE -o json | jq -r '.items[] | select(.spec.containers[0].resources.limits == null) | .meta

    # Check for pending pods due to resource constraints
    echo -e "\n🔒 Pending Pods (Resource Issues):"
    kubectl get pods -n $NAMESPACE --field-selector=status.phase=Pending -o custom-columns=NAME:.metadata.nam

```

```

}

# Function to generate resource recommendations
generate_recommendations() {
    echo -e "\n💡 Resource Optimization Recommendations:"

    # Check for over-provisioned pods (low usage vs high limits)
    echo "1. Check for over-provisioned resources:"
    echo "  kubectl top pods -n $NAMESPACE --sort-by=memory"
    echo "  Compare with limits using: kubectl describe pods -n $NAMESPACE"

    # Check for under-provisioned pods (high usage, low limits)
    echo "2. Monitor for CPU throttling:"
    echo "  kubectl get --raw '/api/v1/nodes/{node-name}/proxy/stats/summary' | jq '.pods[]'"

    # Resource quota recommendations
    echo "3. Consider implementing resource quotas:"
    echo "  kubectl create quota namespace-quota --hard=requests.cpu=4,requests.memory=8Gi,limits.cpu=8,limits.me
}

# Main execution
check_metrics_server
analyze_resources
generate_recommendations

# Create a resource monitoring pod
cat << EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: resource-monitor
  namespace: $NAMESPACE
labels:
  app: resource-monitor
spec:
  containers:
    - name: monitor
      image: busybox
      command: ["sh", "-c"]
      args:
        - |
          while true; do
            echo "$(date): Monitoring resources..."
            # Simple resource usage simulation
EOF

```

```
dd if=/dev/zero of=/tmp/testfile bs=1M count=50 2>/dev/null
rm /tmp/testfile
sleep 60
done
resources:
requests:
  memory: "64Mi"
  cpu: "100m"
limits:
  memory: "128Mi"
  cpu: "200m"
restartPolicy: Always
EOF
```

```
echo "📈 Resource monitor pod created. Check with:"
echo "  kubectl logs -f resource-monitor -n $NAMESPACE"
echo "  kubectl top pod resource-monitor -n $NAMESPACE"
```

## Mini-Project: Automated Resource Right-Sizing

Create a system that analyzes resource usage and suggests optimal resource specifications:

```
bash
```

```
#!/bin/bash
# save as resource-rightsizing.sh
NAMESPACE=${1:-default}
ANALYSIS_PERIOD=${2:-7} # days

echo "⌚ Automated Resource Right-Sizing Analysis"
echo "Namespace: $NAMESPACE"
echo "Analysis Period: $ANALYSIS_PERIOD days"

# Create a comprehensive resource analysis tool
cat << 'EOF' > analyze-resources.py
#!/usr/bin/env python3
import subprocess
import json
import sys
from datetime import datetime, timedelta

def run_kubectl(cmd):
    """Execute kubectl command and return output"""
    try:
        result = subprocess.run(cmd, shell=True, capture_output=True, text=True)
        return result.stdout if result.returncode == 0 else None
    except Exception as e:
        print(f"Error running command: {e}")
        return None

def get_pod_resources(namespace):
    """Get current resource specifications for all pods"""
    cmd = f"kubectl get pods -n {namespace} -o json"
    output = run_kubectl(cmd)
    if not output:
        return []
    pods_data = json.loads(output)
    pod_resources = []

    for pod in pods_data['items']:
        if pod['status']['phase'] != 'Running':
            continue
        pod_info = {
            'name': pod['metadata']['name'],
            'namespace': pod['metadata']['namespace'],
```

```
'containers': []  
}  
  
for container in pod['spec']['containers']:  
    container_info = {  
        'name': container['name'],  
        'requests': container.get('resources', {}).get('requests', {}),  
        'limits': container.get('resources', {}).get('limits', {})  
    }  
    pod_info['containers'].append(container_info)  
  
pod_resources.append(pod_info)  
  
return pod_resources  
  
def get_resource_usage(namespace):  
    """Get current resource usage for pods"""  
    cmd = f"kubectl top pods -n {namespace} --no-headers"  
    output = run_kubectl(cmd)  
    if not output:  
        return {}  
  
    usage = {}  
    for line in output.strip().split('\n'):  
        if line:  
            parts = line.split()  
            if len(parts) >= 3:  
                pod_name = parts[0]  
                cpu_usage = parts[1]  
                memory_usage = parts[2]  
                usage[pod_name] = {  
                    'cpu': cpu_usage,  
                    'memory': memory_usage  
                }  
  
    return usage  
  
def convert_cpu_to_millicores(cpu_str):  
    """Convert CPU string to millicores"""  
    if not cpu_str:  
        return 0  
    if cpu_str.endswith('m'):  
        return int(cpu_str[:-1])  
    else:
```

```
return int(float(cpu_str) * 1000)

def convert_memory_to_mi(memory_str):
    """Convert memory string to Mi"""
    if not memory_str:
        return 0
    if memory_str.endswith('Mi'):
        return int(memory_str[:-2])
    elif memory_str.endswith('Gi'):
        return int(memory_str[:-2]) * 1024
    elif memory_str.endswith('Ki'):
        return int(memory_str[:-2]) // 1024
    else:
        return int(memory_str) // (1024 * 1024)

def analyze_and_recommend(namespace):
    """Main analysis function"""
    print(f"🔍 Analyzing resources for namespace: {namespace}")

    pod_resources = get_pod_resources(namespace)
    current_usage = get_resource_usage(namespace)

    recommendations = []

    for pod in pod_resources:
        pod_name = pod['name']
        usage = current_usage.get(pod_name, {})

        if not usage:
            continue

        for container in pod['containers']:
            container_name = container['name']
            requests = container['requests']
            limits = container['limits']

            # Current usage
            current_cpu = convert_cpu_to_millicores(usage.get('cpu', '0'))
            current_memory = convert_memory_to_mi(usage.get('memory', '0'))

            # Current limits
            limit_cpu = convert_cpu_to_millicores(limits.get('cpu', '0'))
            limit_memory = convert_memory_to_mi(limits.get('memory', '0'))
```

```

# Current requests
request_cpu = convert_cpu_to_millicores(requests.get('cpu', '0'))
request_memory = convert_memory_to_mi(requests.get('memory', '0'))

# Calculate recommendations (usage * 1.2 for headroom)
recommended_cpu = max(current_cpu * 1.2, 100) # minimum 100m
recommended_memory = max(current_memory * 1.2, 64) # minimum 64Mi

recommendation = {
    'pod': pod_name,
    'container': container_name,
    'current_usage': {
        'cpu': f'{current_cpu}m',
        'memory': f'{current_memory}Mi'
    },
    'current_limits': {
        'cpu': f'{limit_cpu}m' if limit_cpu else "none",
        'memory': f'{limit_memory}Mi' if limit_memory else "none"
    },
    'recommended': {
        'cpu_request': f'{int(recommended_cpu)}m',
        'cpu_limit': f'{int(recommended_cpu * 1.5)}m',
        'memory_request': f'{int(recommended_memory)}Mi',
        'memory_limit': f'{int(recommended_memory * 1.5)}Mi'
    }
}

recommendations.append(recommendation)

return recommendations

def generate_yaml(recommendations):
    """Generate YAML patches for resource recommendations"""
    print("\nGenerated Resource Patches:")

    for rec in recommendations:
        print(f"\n# Patch for {rec['pod']} - {rec['container']}")
        print(f"kubectl patch pod {rec['pod']} -p '{")
        print(f'    "spec": {{"')
        print(f'        "containers": [{"')
        print(f'            "name": "{rec["container"]}"},')
        print(f'            "resources": {{"')
        print(f'                "requests": {{"')
        print(f'                    "cpu": "{rec["recommended"]["cpu_request"]}"')

```

```

print(f'      "memory": "{rec["recommended"]["memory_request"]}"')
print(f'    }},')
print(f'    "limits": {"')
print(f'      "cpu": "{rec["recommended"]["cpu_limit"]}","')
print(f'      "memory": "{rec["recommended"]["memory_limit"]}"}')
print(f'    }')
print(f'  }]')
print(f'  }}')
print(f"  }")

```

if \_\_name\_\_ == "\_\_main\_\_":
 namespace = sys.argv[1] if len(sys.argv) > 1 else "default"

recommendations = analyze\_and\_recommend(namespace)

print("\n

generate\_yaml(recommendations)

EOF

[chmod](#) +x analyze-resources.py

```

# Run the analysis
echo "🚀 Running resource analysis.."
python3 analyze-resources.py $NAMESPACE

# Create deployment with various resource patterns for testing
cat << EOF | kubectl apply -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: resource-test-app
  namespace: $NAMESPACE
spec:

```

```
replicas: 3
selector:
  matchLabels:
    app: resource-test
template:
  metadata:
    labels:
      app: resource-test
spec:
  containers:
    - name: web
      image: nginx:alpine
      resources:
        requests:
          memory: "64Mi"
          cpu: "100m"
        limits:
          memory: "128Mi"
          cpu: "200m"
    - name: sidecar
      image: busybox
      command: ["sh", "-c", "while true; do echo 'sidecar running'; sleep 30; done"]
      resources:
        requests:
          memory: "32Mi"
          cpu: "50m"
        limits:
          memory: "64Mi"
          cpu: "100m"
EOF
```

```
echo "✅ Resource right-sizing analysis complete!"
echo "📈 Monitor resource usage over time with:"
echo "  watch kubectl top pods -n $NAMESPACE"
echo "⚡ Cleanup:"
echo "  kubectl delete deployment resource-test-app -n $NAMESPACE"
echo "  rm analyze-resources.py"
```

## 19. Helm Basics

Helm is the package manager for Kubernetes, allowing you to define, install, and upgrade complex Kubernetes applications using templates and charts.

## 19.1 Installing and Using Helm

```
bash
```

```
# Install Helm (if not already installed)
curl https://get.helm.sh/helm-v3.12.0-linux-amd64.tar.gz | tar xz
sudo mv linux-amd64/helm /usr/local/bin/

# Verify Helm installation
helm version

# Add popular Helm repositories
helm repo add bitnami https://charts.bitnami.com/bitnami
helm repo add stable https://charts.helm.sh/stable
helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx

# Update repository information
helm repo update

# Search for charts
helm search repo nginx
helm search repo database

# Install a chart
helm install my-nginx bitnami/nginx
# Creates a release named "my-nginx" using the bitnami/nginx chart

# List installed releases
helm list
helm list --all-namespaces

# Get information about a release
helm status my-nginx
helm get values my-nginx
helm get manifest my-nginx

# Upgrade a release
helm upgrade my-nginx bitnami/nginx --set service.type=NodePort

# Rollback a release
helm rollback my-nginx 1

# Uninstall a release
helm uninstall my-nginx
```

## Example: Creating Your First Helm Chart

Let's create a custom Helm chart for a web application:

```
bash
```

```
#!/bin/bash
# save as create-helm-chart.sh
CHART_NAME=${1:-webapp}

echo "📦 Creating Helm Chart: $CHART_NAME"

# Create a new Helm chart
helm create $CHART_NAME

echo "📁 Chart structure created:"
tree $CHART_NAME

# Customize the chart values
cat << EOF > $CHART_NAME/values.yaml
# Default values for $CHART_NAME
replicaCount: 2

image:
  repository: nginx
  pullPolicy: IfNotPresent
  tag: "1.21"

nameOverride: ""
fullnameOverride: ""

service:
  type: ClusterIP
  port: 80

ingress:
  enabled: true
  className: ""
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/rewrite-target: /
  hosts:
    - host: $CHART_NAME.local
      paths:
        - path: /
          pathType: Prefix
  tls: []

resources:
```

```
limits:
  cpu: 500m
  memory: 512Mi
requests:
  cpu: 250m
  memory: 256Mi

autoscaling:
  enabled: true
  minReplicas: 2
  maxReplicas: 10
  targetCPUUtilizationPercentage: 80

nodeSelector: {}
tolerations: []
affinity: {}

# Custom application configuration
app:
  environment: production
  debug: false
database:
  host: db.example.com
  port: 5432
  name: myapp
EOF
```

```
# Create a ConfigMap template for application configuration
cat << 'EOF' > ${CHART_NAME}/templates/configmap.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: {{ include "webapp.fullname" . }}-config
labels:
  {{- include "webapp.labels" . | nindent 4 }}
data:
  app.properties: |
    environment={{ .Values.app.environment }}
    debug={{ .Values.app.debug }}
    database.host={{ .Values.app.database.host }}
    database.port={{ .Values.app.database.port | quote }}
    database.name={{ .Values.app.database.name }}
EOF
```

```
# Update deployment template to use ConfigMap
cat << 'EOF' > ${CHART_NAME}/templates/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ include "webapp.fullname" . }}
  labels:
    {{- include "webapp.labels" . | nindent 4 }}
spec:
  {{- if not .Values.autoscaling.enabled --}}
  replicas: {{ .Values.replicaCount }}
  {{- end --}}
selector:
  matchLabels:
    {{- include "webapp.selectorLabels" . | nindent 6 --}}
template:
  metadata:
    annotations:
      checksum/config: {{ include (print $.Template.BasePath "/configmap.yaml") . | sha256sum }}
    labels:
      {{- include "webapp.selectorLabels" . | nindent 8 --}}
spec:
  containers:
    - name: {{ .Chart.Name }}
      image: "{{ .Values.image.repository }}:{{ .Values.image.tag | default .Chart.AppVersion }}"
      imagePullPolicy: {{ .Values.image.pullPolicy }}
      ports:
        - name: http
          containerPort: 80
          protocol: TCP
      livenessProbe:
        httpGet:
          path: /
          port: http
      initialDelaySeconds: 30
      periodSeconds: 10
      readinessProbe:
        httpGet:
          path: /
          port: http
      initialDelaySeconds: 5
      periodSeconds: 5
  resources:
    {{- toYaml .Values.resources | nindent 12 }}
```

```

volumeMounts:
  - name: config
    mountPath: /etc/config
    readOnly: true
  env:
    - name: ENVIRONMENT
      value: {{ .Values.app.environment | quote }}
volumes:
  - name: config
    configMap:
      name: {{ include "webapp.fullname" . }}-config
      {{- with .Values.nodeSelector --}}
nodeSelector:
  {{- toYaml . | nindent 8 --}}
{{- end --}}
{{- with .Values.affinity --}}
affinity:
  {{- toYaml . | nindent 8 --}}
{{- end --}}
{{- with .Values.tolerations --}}
tolerations:
  {{- toYaml . | nindent 8 --}}
{{- end --}}
EOF

```

`echo "✅ Helm chart created successfully!"`

*# Validate the chart*

`echo "🔍 Validating chart..."`

`helm lint $CHART_NAME`

*# Test chart rendering*

`echo "📝 Testing chart templates..."`

`helm template $CHART_NAME ./${CHART_NAME}`

*# Install the chart*

`echo "🚀 Installing chart..."`

`helm install $CHART_NAME-release ./${CHART_NAME}`

`echo "📊 Chart deployed! Check status with:"`

`echo " helm status $CHART_NAME-release"`

`echo " kubectl get all -l app.kubernetes.io/instance=$CHART_NAME-release"`

`echo "🧹 Cleanup:"`

```
echo " helm uninstall $CHART_NAME-release"
echo " rm -rf $CHART_NAME"
```

## Demo: Multi-Environment Deployment with Helm

Create a system for deploying applications across different environments:

bash

```
#!/bin/bash
# save as multi-env-helm-demo.sh
APP_NAME="myapp"
ENVIRONMENTS=(“development” “staging” “production”)

echo “🌐 Multi-Environment Helm Deployment Demo”

# Create base chart
helm create $APP_NAME

# Create environment-specific values files
create_env_values() {
    local env=$1
    local replicas cpu_limit memory_limit ingress_host

    case $env in
        “development”)
            replicas=1
            cpu_limit=“200m”
            memory_limit=“256Mi”
            ingress_host=“dev-$APP_NAME.local”
            ;;
        “staging”)
            replicas=2
            cpu_limit=“500m”
            memory_limit=“512Mi”
            ingress_host=“staging-$APP_NAME.local”
            ;;
        “production”)
            replicas=3
            cpu_limit=“1000m”
            memory_limit=“1Gi”
            ingress_host=“$APP_NAME.example.com”
            ;;
    esac

    cat << EOF > $APP_NAME/values-$env.yaml
# Values for $env environment
replicaCount: $replicas

image:
repository: nginx
tag: “1.21”
EOF
done
}
```

```
pullPolicy: IfNotPresent

service:
  type: ClusterIP
  port: 80

ingress:
  enabled: true
  className: "nginx"
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
  hosts:
    - host: $ingress_host
      paths:
        - path: /
  pathType: Prefix

resources:
  limits:
    cpu: $cpu_limit
    memory: $memory_limit
  requests:
    cpu: $(echo $cpu_limit | sed 's/[0-9]*m//' | awk '{print $1/2"m"}' | sed 's/\..*m/m/')
    memory: $(echo $memory_limit | sed 's/[0-9]*Mi//' | awk '{print $1/2"Mi"}' | sed 's/\..*Mi/Mi/')

autoscaling:
  enabled: ${([ "$env" = "production" ] && echo "true" || echo "false")}
  minReplicas: $replicas
  maxReplicas: $($replicas * 3)
  targetCPUUtilizationPercentage: 80

# Environment-specific configuration
env:
  name: $env
  debug: ${([ "$env" = "development" ] && echo "true" || echo "false")}
  logLevel: ${([ "$env" = "production" ] && echo "warn" || echo "info")}

EOF
}

# Create environment-specific values
for env in "${ENVIRONMENTS[@]}"; do
  echo "📝 Creating values for $env environment..."
  create_env_values $env
done
```

```

# Create namespaces for each environment
for env in "${ENVIRONMENTS[@]}"; do
    echo "🏗️ Creating namespace for $env..."
    kubectl create namespace $env --dry-run=client -o yaml | kubectl apply -f -
done

# Deploy to each environment
deploy_to_env() {
    local env=$1
    echo "🚀 Deploying $APP_NAME to $env environment..."

    helm upgrade --install $APP_NAME-$env ./${APP_NAME} \
        --namespace $env \
        --values ./${APP_NAME}/values-$env.yaml \
        --wait \
        --timeout 300s

    echo "✅ Deployment to $env complete!"
}

# Deploy to all environments
for env in "${ENVIRONMENTS[@]}"; do
    deploy_to_env $env
done

echo "📊 Deployment Summary:"
for env in "${ENVIRONMENTS[@]}"; do
    echo "--- $env Environment ---"
    helm status ${APP_NAME}-$env -n $env
    kubectl get pods -n $env -l app.kubernetes.io/instance=${APP_NAME}-$env
done

echo "🔍 Check deployments:"
echo " helm list --all-namespaces"
echo " kubectl get pods --all-namespaces -l app.kubernetes.io/name=${APP_NAME}"

echo "🧹 Cleanup all environments:"
for env in "${ENVIRONMENTS[@]}"; do
    echo " helm uninstall ${APP_NAME}-$env -n $env"
done
echo " kubectl delete namespace ${ENVIRONMENTS[*]}"
echo " rm -rf ${APP_NAME}"

```

## Mini-Project: Helm Chart Repository and CI/CD Pipeline

Create a complete Helm chart repository with automated testing and deployment:

```
bash
```

```
#!/bin/bash
# save as helm-cicd-pipeline.sh
REPO_NAME="myapp-charts"
CHART_NAME="webapp"

echo "🏗️ Helm Chart Repository and CI/CD Pipeline Setup"

# Create chart repository structure
mkdir -p $REPO_NAME/{charts,docs,scripts,tests}
cd $REPO_NAME

# Initialize git repository
git init
cat << EOF > .gitignore
*.tgz
.DS_Store
Thumbs.db
EOF

# Create base chart
helm create charts/$CHART_NAME

# Create chart testing configuration
cat << EOF > charts/$CHART_NAME/ci/test-values.yaml
# Test values for CI pipeline
replicaCount: 1

image:
repository: nginx
tag: "alpine"
pullPolicy: IfNotPresent

service:
type: ClusterIP
port: 80

resources:
limits:
cpu: 100m
memory: 128Mi
requests:
cpu: 50m
memory: 64Mi
```

```
# Minimal configuration for testing
autoscaling:
  enabled: false

ingress:
  enabled: false
EOF

# Create comprehensive test suite
cat << 'EOF' > tests/chart-test.sh
#!/bin/bash
set -e

CHART_PATH="charts/webapp"
RELEASE_NAME="test-release"
NAMESPACE="chart-testing"

echo "🧪 Running Helm Chart Tests"

# Create testing namespace
kubectl create namespace $NAMESPACE --dry-run=client -o yaml | kubectl apply -f -

# Lint the chart
echo "📋 Linting chart..."
helm lint $CHART_PATH

# Test template rendering
echo "🎨 Testing template rendering..."
helm template $RELEASE_NAME $CHART_PATH --values $CHART_PATH/ci/test-values.yaml > /tmp/rendered-templates.yaml

# Validate generated YAML
echo "✅ Validating generated YAML..."
kubectl apply --dry-run=client -f /tmp/rendered-templates.yaml

# Install chart for integration testing
echo "🚀 Installing chart for testing..."
helm install $RELEASE_NAME $CHART_PATH \
  --namespace $NAMESPACE \
  --values $CHART_PATH/ci/test-values.yaml \
  --wait \
  --timeout 300s

# Run tests
```

```

echo "🔍 Running integration tests..."

# Test deployment status
if kubectl get deployment $RELEASE_NAME-webapp -n $NAMESPACE >/dev/null 2>&1; then
    echo "✅ Deployment created successfully"
else
    echo "❌ Deployment not found"
    exit 1
fi

# Test service creation
if kubectl get service $RELEASE_NAME-webapp -n $NAMESPACE >/dev/null 2>&1; then
    echo "✅ Service created successfully"
else
    echo "❌ Service not found"
    exit 1
fi

# Test pod readiness
echo "⌚ Waiting for pods to be ready..."
kubectl wait --for=condition=ready pod -l app.kubernetes.io/instance=$RELEASE_NAME -n $NAMESPACE --timeout=1m
echo "✅ Pods are ready"

# Test HTTP connectivity
echo "🌐 Testing HTTP connectivity..."
kubectl run test-pod --image=busybox --rm -it --restart=Never -n $NAMESPACE -- \
    wget -q --spider http://$RELEASE_NAME-webapp/
echo "✅ HTTP connectivity test passed"

# Cleanup
echo "🧹 Cleaning up test resources..."
helm uninstall $RELEASE_NAME -n $NAMESPACE
kubectl delete namespace $NAMESPACE

echo "🎉 All tests passed!"
EOF

chmod +x tests/chart-test.sh

# Create release automation script
cat << 'EOF' > scripts/release.sh
#!/bin/bash
set -e

```

```
CHART_PATH="charts/webapp"
VERSION=$1

if [ -z "$VERSION" ]; then
    echo "Usage: $0 <version>"
    echo "Example: $0 1.0.0"
    exit 1
fi

echo "📦 Creating Helm Chart Release v$VERSION"

# Update chart version
sed -i "s/version:.*version: $VERSION/" $CHART_PATH/Chart.yaml
sed -i "s/appVersion:.*appVersion: \"$VERSION\"" $CHART_PATH/Chart.yaml

# Run tests
echo "🧪 Running tests before release..."
./tests/chart-test.sh

# Package chart
echo "📦 Packaging chart..."
helm package $CHART_PATH --destination ./docs/

# Update repository index
echo "📋 Updating repository index..."
helm repo index ./docs/ --url https://mycompany.github.io/myapp-charts/

# Create GitHub Pages index
cat << 'HTML' > docs/index.html
<!DOCTYPE html>
<html>
<head>
    <title>MyApp Helm Charts</title>
    <style>
        body { font-family: Arial, sans-serif; margin: 40px; }
        .chart { background: #f5f5f5; padding: 20px; margin: 20px 0; border-radius: 5px; }
        .version { color: #666; }
    </style>
</head>
<body>
    <h1>🚢 MyApp Helm Chart Repository</h1>
    <p>Add this repository to Helm:</p>
    <pre><code>helm repo add myapp-charts https://mycompany.github.io/myapp-charts/</code></pre>
```

```
<div class="chart">
  <h2>webapp</h2>
  <p class="version">Latest Version: VERSION_PLACEHOLDER</p>
  <p>A production-ready web application chart with configurable environments.</p>
  <pre><code>helm install my-webapp myapp-charts/webapp</code></pre>
</div>
</body>
</html>
```

HTML

```
# Replace version placeholder
sed -i "s/VERSION_PLACEHOLDER/$VERSION/" docs/index.html
```

```
# Git operations
git add .
git commit -m "Release v$VERSION"
git tag -a "v$VERSION" -m "Release version $VERSION"
```

```
echo "✅ Release v$VERSION created successfully!"
echo "👉 Push to repository with:"
echo " git push origin main"
echo " git push origin v$VERSION"
EOF
```

```
chmod +x scripts/release.sh
```

```
# Create GitHub Actions workflow
mkdir -p .github/workflows
cat << 'EOF' > .github/workflows/ci.yml
name: Chart CI/CD
```

```
on:
```

```
  push:
    branches: [ main ]
  pull_request:
    branches: [ main ]
```

```
jobs:
```

```
  test:
```

```
    runs-on: ubuntu-latest
```

```
    steps:
```

```
      - uses: actions/checkout@v3
```

```
      - name: Set up Helm
```

```
uses: azure/setup-helm@v3
```

```
with:
```

```
version: '3.12.0'
```

```
- name: Set up kubectl
```

```
uses: azure/setup-kubectl@v3
```

```
- name: Create k8s Kind Cluster
```

```
uses: helm/kind-action@v1.4.0
```

```
- name: Run chart tests
```

```
run: ./tests/chart-test.sh
```

```
release:
```

```
needs: test
```

```
runs-on: ubuntu-latest
```

```
if: github.ref == 'refs/heads/main'
```

```
steps:
```

```
- uses: actions/checkout@v3
```

```
with:
```

```
token: ${{ secrets.GITHUB_TOKEN }}
```

```
- name: Set up Helm
```

```
uses: azure/setup-helm@v3
```

```
with:
```

```
version: '3.12.0'
```

```
- name: Configure Git
```

```
run: |
```

```
git config user.name "$GITHUB_ACTOR"
```

```
git config user.email "$GITHUB_ACTOR@users.noreply.github.com"
```

```
- name: Release Chart
```

```
run: |
```

```
VERSION=$(date +%Y.%m.%d-%H%M%S)
```

```
./scripts/release.sh $VERSION
```

```
- name: Deploy to GitHub Pages
```

```
uses: peaceiris/actions-gh-pages@v3
```

```
with:
```

```
github_token: ${{ secrets.GITHUB_TOKEN }}
```

```
publish_dir: ./docs
```

```
EOF
```

```
# Create comprehensive documentation
cat << EOF > README.md
# 🚀 MyApp Helm Charts
```

This repository contains Helm charts for MyApp deployments.

## ## Usage

Add the repository:

```
\`\`\`bash
helm repo add myapp-charts https://mycompany.github.io/myapp-charts/
helm repo update
\`\`\`
```

Install a chart:

```
\`\`\`bash
helm install my-webapp myapp-charts/webapp
\`\`\`
```

## ## Charts

### ### webapp

A production-ready web application chart with:

- Multi-environment support
- Horizontal Pod Autoscaling
- Ingress configuration
- Resource management
- Health checks

## ## Development

### ### Testing

Run the test suite:

```
\`\`\`bash
./tests/chart-test.sh
\`\`\`
```

### ### Releasing

Create a new release:

```
\`\`\`bash
./scripts/release.sh 1.0.0
\`\`\`
```

## ## Contributing

1. Make changes to charts
  2. Run tests locally
  3. Create pull request
  4. Tests run automatically
  5. Merge triggers automatic release
- EOF

```
echo "✅ Helm Chart Repository Setup Complete!"  
echo ""  
echo "📁 Repository Structure:"  
find . -type f -name "*.sh" -o -name "*.yaml" -o -name "*.yml" -o -name "*.md" | head -20  
  
echo ""  
echo "🚀 Next Steps:"  
echo "1. Initialize git: git add . && git commit -m 'Initial commit'"  
echo "2. Run tests: ./tests/chart-test.sh"  
echo "3. Create release: ./scripts/release.sh 1.0.0"  
echo "4. Push to GitHub and enable Pages"  
  
cd ..
```

## 20. Configuration and Contexts

Managing multiple Kubernetes clusters and contexts is essential for development workflows and multi-environment deployments.

### 20.1 Managing Contexts and Configurations

```
bash
```

```
# View current configuration
kubectl config view
# Shows merged configuration from all kubeconfig files

# List available contexts
kubectl config get-contexts
# Shows all configured clusters, users, and contexts

# Display current context
kubectl config current-context
# Shows which cluster you're currently connected to

# Switch context
kubectl config use-context minikube
kubectl config use-context production-cluster

# Set default namespace for current context
kubectl config set-context --current --namespace=development
# All kubectl commands will use this namespace by default

# Create a new context
kubectl config set-context dev-context \
  --cluster=development-cluster \
  --user=dev-user \
  --namespace=development

# Rename context
kubectl config rename-context old-name new-name

# Delete context
kubectl config delete-context unused-context
```

## Example: Multi-Cluster Management Setup

Create a comprehensive multi-cluster management system:

```
bash
```

```

#!/bin/bash
# save as multi-cluster-setup.sh
echo "🌐 Multi-Cluster Kubernetes Management Setup"

# Create directory for cluster configurations
mkdir -p ~/kube/clusters
KUBECONFIG_DIR="$HOME/.kube/clusters"

# Function to create mock cluster configurations
create_cluster_config() {
    local cluster_name=$1
    local environment=$2
    local endpoint=$3

    cat << EOF > $KUBECONFIG_DIR/$cluster_name-config
apiVersion: v1
kind: Config
clusters:
- cluster:
    server: $endpoint
    certificate-authority-data: LS0tLS1CRUdJTi... # Mock cert data
    name: $cluster_name
contexts:
- context:
    cluster: $cluster_name
    user: $cluster_name-user
    namespace: default
    name: $cluster_name
current-context: $cluster_name
users:
- name: $cluster_name-user
user:
    token: mock-token-for-$cluster_name
EOF

    echo "✅ Created config for $cluster_name ($environment)"
}

# Create configurations for different environments
create_cluster_config "dev-cluster" "development" "https://dev-k8s-api.company.com"
create_cluster_config "staging-cluster" "staging" "https://staging-k8s-api.company.com"
create_cluster_config "prod-cluster" "production" "https://prod-k8s-api.company.com"

```

```
# Create context management script
cat << 'EOF' > ~/kube/kube-context-manager.sh
#!/bin/bash

# Function to display available contexts
show_contexts() {
    echo "📋 Available Kubernetes Contexts:"
    kubectl config get-contexts
}

# Function to switch context with safety checks
switch_context() {
    local target_context=$1

    if [ -z "$target_context" ]; then
        echo "✖ Please specify a context name"
        show_contexts
        return 1
    fi

    # Check if context exists
    if ! kubectl config get-contexts -o name | grep -q "^$target_context$"; then
        echo "✖ Context '$target_context' not found"
        show_contexts
        return 1
    fi

    # Safety check for production
    if [[ "$target_context" == *"prod"* ]]; then
        echo "⚠ You're about to switch to PRODUCTION context!"
        read -p "Are you sure? (yes/no): " confirmation
        if [ "$confirmation" != "yes" ]; then
            echo "✖ Context switch cancelled"
            return 1
        fi
    fi

    # Switch context
    kubectl config use-context $target_context

    # Display current context info
    echo "✓ Switched to context: $target_context"
    echo "🔍 Current cluster info:"
    kubectl cluster-info --context=$target_context
```

```
# Set PS1 to show current context
export PS1="[k8s:$target_context] \u@\h:\w\$ "
}

# Function to create context alias
create_alias() {
    local alias_name=$1
    local context_name=$2

    if [ -z "$alias_name" ] || [ -z "$context_name" ]; then
        echo "Usage: create_alias <alias> <context>"
        return 1
    fi

    echo "alias k-$alias_name='kubectl config use-context $context_name'" >> ~/.bashrc
    echo "✓ Created alias 'k-$alias_name' for context '$context_name'"
    echo "Run 'source ~/.bashrc' to activate"
}

# Function to backup current kubeconfig
backup_config() {
    local backup_name="kubeconfig-backup-$(date +%Y%m%d-%H%M%S)"
    cp ~/.kube/config ~/.kube/$backup_name
    echo "✓ Kubeconfig backed up to ~/.kube/$backup_name"
}

# Function to merge kubeconfig files
merge_configs() {
    echo "🔄 Merging kubeconfig files from ~/.kube/clusters/"

    # Backup current config
    backup_config

    # Merge all configs
    export KUBECONFIG=~/kube/config:$(find ~/kube/clusters -name "*-config" | tr '\n' ':')
    kubectl config view --flatten > ~/kube/config-merged
    mv ~/kube/config-merged ~/kube/config

    echo "✓ Configurations merged successfully"
    show_contexts
}

# Main command handler
```

```

case "$1" in
"list"|"ls")
    show_contexts
;;
"switch"|"use")
    switch_context $2
;;
"alias")
    create_alias $2 $3
;;
"backup")
    backup_config
;;
"merge")
    merge_configs
;;
"current")
    echo "Current context: $(kubectl config current-context)"
    kubectl config get-contexts $(kubectl config current-context)
;;
*)
    echo "🚀 Kubernetes Context Manager"
    echo "Usage: $0 {list|switch|alias|backup|merge|current}"
    echo ""
    echo "Commands:"
    echo " list           - Show available contexts"
    echo " switch <context>   - Switch to specified context"
    echo " alias <name> <context> - Create alias for context"
    echo " backup          - Backup current kubeconfig"
    echo " merge           - Merge configs from clusters directory"
    echo " current         - Show current context"
;;
esac
EOF

```

[chmod +x ./kube/context-manager.sh](#)

```

# Create convenient aliases
echo "# Kubernetes context management aliases" >> ~/.bashrc
echo "alias kctx='~/kube/context-manager.sh'" >> ~/.bashrc
echo "alias k-dev='kubectl config use-context dev-cluster'" >> ~/.bashrc
echo "alias k-staging='kubectl config use-context staging-cluster'" >> ~/.bashrc
echo "alias k-prod='kubectl config use-context prod-cluster'" >> ~/.bashrc

```

```
echo "✅ Multi-cluster setup complete!"  
echo ""  
echo "👉 Usage:"  
echo " kctx list      # Show all contexts"  
echo " kctx switch dev-cluster # Switch to dev cluster"  
echo " k-dev          # Quick switch to dev"  
echo " k-staging       # Quick switch to staging"  
echo " k-prod          # Quick switch to production"  
echo ""  
echo "📝 Run 'source ~/.bashrc' to activate aliases"
```

## Demo: Environment-Specific Configuration Management

Create a system for managing different configurations per environment:

bash

```
#!/bin/bash
# save as env-config-demo.sh
echo "🔧 Environment-Specific Configuration Management Demo"

# Create environment configurations
ENVIRONMENTS=(“development” “staging” “production”)

create_env_config() {
local env=$1
mkdir -p configs/$env

# Environment-specific kubectl config
cat << EOF > configs/$env/kubeconfig
apiVersion: v1
kind: Config
clusters:
- cluster:
  server: https://$env-api.example.com
  name: $env-cluster
contexts:
- context:
  cluster: $env-cluster
  user: $env-user
  namespace: $env
  name: $env
current-context: $env
users:
- name: $env-user
user:
  token: $env-token-12345
EOF

# Environment-specific resource quotas
case $env in
  “development”)
    cat << EOF > configs/$env/resource-quota.yaml
apiVersion: v1
kind: ResourceQuota
metadata:
  name: $env-quota
  namespace: $env
spec:
  hard:
```

```
requests.cpu: "4"
requests.memory: 8Gi
limits.cpu: "8"
limits.memory: 16Gi
pods: "20"
services: "10"

EOF
;;
"staging")
cat << EOF > configs/$env/resource-quota.yaml
apiVersion: v1
kind: ResourceQuota
metadata:
  name: $env-quota
  namespace: $env
spec:
  hard:
    requests.cpu: "8"
    requests.memory: 16Gi
    limits.cpu: "16"
    limits.memory: 32Gi
    pods: "50"
    services: "20"

EOF
;;
"production")
cat << EOF > configs/$env/resource-quota.yaml
apiVersion: v1
kind: ResourceQuota
metadata:
  name: $env-quota
  namespace: $env
spec:
  hard:
    requests.cpu: "20"
    requests.memory: 40Gi
    limits.cpu: "40"
    limits.memory: 80Gi
    pods: "100"
    services: "50"

EOF
;;
esac
```

```
# Environment-specific network policies
cat << EOF > configs/$env/network-policy.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: $env-network-policy
  namespace: $env
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            environment: $env
  egress:
    - to: []
EOF

echo "✓ Created configuration for $env environment"
}
```

```
# Create configurations for all environments
for env in "${ENVIRONMENTS[@]}"; do
  create_env_config $env
done
```

```
# Create environment switcher script
cat << 'EOF' > switch-env.sh
#!/bin/bash
ENV=$1

if [ -z "$ENV" ]; then
  echo "Usage: $0 {development|staging|production}"
  exit 1
fi

if [ ! -d "configs/$ENV" ]; then
  echo "✗ Environment '$ENV' not found"
  exit 1
fi
```

```
echo "⚡️ Switching to $ENV environment...""

# Backup current kubeconfig
cp ~/.kube/config ~/.kube/config.backup

# Set environment-specific kubeconfig
export KUBECONFIG=~/.kube/config:configs/$ENV/kubeconfig
kubectl config view --flatten > ~/.kube/config

# Switch to environment context
kubectl config use-context $ENV

# Create namespace if it doesn't exist
kubectl create namespace $ENV --dry-run=client -o yaml | kubectl apply -f -

# Apply environment-specific configurations
echo "📝 Applying environment configurations..."
kubectl apply -f configs/$ENV/resource-quota.yaml
kubectl apply -f configs/$ENV/network-policy.yaml

# Set default namespace
kubectl config set-context --current --namespace=$ENV

echo "✅ Successfully switched to $ENV environment"
echo "🌐 Current context: $(kubectl config current-context)"
echo "👉 Default namespace: $(kubectl config view --minify | grep namespace | awk '{print $2}')"

# Update shell prompt
export PS1="[$ENV-k8s] \u@\:~\$ "
EOF

chmod +x switch-env.sh

echo "📝 Environment configurations created:"
find configs -name "*.yaml" | head -10

echo """
echo "🚀 Usage:"
echo " ./switch-env.sh development # Switch to dev environment"
echo " ./switch-env.sh staging    # Switch to staging environment"
echo " ./switch-env.sh production # Switch to production environment"

echo """
echo "🔍 Verify configuration:"
```

```
echo " kubectl config get-contexts"
echo " kubectl get resourcequota --all-namespaces"
```

## Mini-Project: Advanced Kubeconfig Management System

Create a comprehensive system for managing complex multi-cluster, multi-user scenarios:

bash

```
#!/bin/bash
# save as advanced-kubeconfig-manager.sh
echo "🌐 Advanced Kubeconfig Management System"

# Create directory structure
mkdir -p ~/kube/{contexts,users,clusters,backups,profiles}

# Create user management functions
create_user_profile() {
    local username=$1
    local email=$2
    local role=$3

    cat << EOF > ~/.kube/users/$username.yaml
apiVersion: v1
kind: Config
preferences: {}
users:
- name: $username
  user:
    client-certificate-data: $(echo "cert-data-for-$username" | base64)
    client-key-data: $(echo "key-data-for-$username" | base64)
  metadata:
    email: $email
    role: $role
    created: $(date -Iseconds)
EOF

    echo "👤 Created user profile for $username ($role)"
}

# Create cluster profiles
```

```
create_cluster_profile() {
    local cluster_name=$1
    local environment=$2
    local endpoint=$3
    local region=$4

    cat << EOF > ~/.kube/clusters/$cluster_name.yaml
apiVersion: v1
kind: Config
preferences: {}
clusters:
```

```
- cluster:
    certificate-authority-data: $(echo "ca-data-for-$cluster_name" | base64)
    server: $endpoint
    name: $cluster_name
metadata:
    environment: $environment
    region: $region
    created: $(date -Iseconds)
EOF

echo "💡 Created cluster profile for $cluster_name ($environment - $region)"
}
```

```
# Create context templates
create_context_template() {
    local context_name=$1
    local cluster=$2
    local user=$3
    local namespace=$4
    local description=$5
```

```
    cat << EOF > ~/.kube-contexts/$context_name.yaml
apiVersion: v1
kind: Config
preferences: {}
contexts:
- context:
    cluster: $cluster
    user: $user
    namespace: $namespace
    name: $context_name
current-context: $context_name
metadata:
    description: "$description"
    created: $(date -Iseconds)
EOF
```

```
echo "🔗 Created context template for $context_name"
}
```

```
# Advanced kubeconfig manager script
cat << 'EOF' > ~/.kube/advanced-kubeconfig-manager.sh
#!/bin/bash
```

```

KUBE_DIR="$HOME/.kube"
BACKUP_DIR="$KUBE_DIR/backups"
PROFILES_DIR="$KUBE_DIR/profiles"

# Function to create comprehensive backup
create_backup() {
    local backup_name="full-backup-$(date +%Y%m%d-%H%M%S)"
    local backup_path="$BACKUP_DIR/$backup_name"

    mkdir -p "$backup_path"

    # Backup current config
    cp "$KUBE_DIR/config" "$backup_path/config" 2>/dev/null || true

    # Backup all profiles
    cp -r "$KUBE_DIR"/{contexts,users,clusters,profiles} "$backup_path/" 2>/dev/null || true

    # Create backup manifest
    cat << MANIFEST > "$backup_path/manifest.yaml"
    backup_name: $backup_name
    created: $(date -Iseconds)
    kubectl_version: $(kubectl version --client --short 2>/dev/null || echo "unknown")
    contexts_count: $(kubectl config get-contexts --no-headers 2>/dev/null | wc -l || echo "0")
    current_context: $(kubectl config current-context 2>/dev/null || echo "none")
    MANIFEST

    echo "💾 Backup created: $backup_path"
}

# Function to build kubeconfig from profiles
build_config() {
    local profile_name=${1:-"default"}
    local temp_config="/tmp/kubeconfig-build-$"

    echo "🔨 Building kubeconfig profile: $profile_name"

    # Start with empty config
    cat << CONFIG > "$temp_config"
    apiVersion: v1
    kind: Config
    clusters: []
    contexts: []
    users: []
    preferences: {}

```

```
current-context: ""

CONFIG

# Read profile configuration
local profile_file="$PROFILES_DIR/$profile_name.yaml"
if [ ! -f "$profile_file" ]; then
    echo "✖ Profile '$profile_name' not found"
    return 1
fi

# Source profile settings
source "$profile_file"

# Merge configurations
export KUBECONFIG="$temp_config"

# Add clusters
for cluster_file in $CLUSTERS; do
    if [ -f "$KUBE_DIR/clusters/$cluster_file.yaml" ]; then
        export KUBECONFIG="$KUBECONFIG:$KUBE_DIR/clusters/$cluster_file.yaml"
    fi
done

# Add users
for user_file in $USERS; do
    if [ -f "$KUBE_DIR/users/$user_file.yaml" ]; then
        export KUBECONFIG="$KUBECONFIG:$KUBE_DIR/users/$user_file.yaml"
    fi
done

# Add contexts
for context_file in $CONTEXTS; do
    if [ -f "$KUBE_DIR/contexts/$context_file.yaml" ]; then
        export KUBECONFIG="$KUBECONFIG:$KUBE_DIR/contexts/$context_file.yaml"
    fi
done

# Flatten and save
kubectl config view --flatten > "$KUBE_DIR/config-$profile_name"

# Set default context if specified
if [ ! -z "$DEFAULT_CONTEXT" ]; then
    KUBECONFIG="$KUBE_DIR/config-$profile_name" kubectl config use-context "$DEFAULT_CONTEXT"
fi
```

```
echo "✅ Profile '$profile_name' built successfully"
echo "📄 Config saved to: $KUBE_DIR/config-$profile_name"

rm "$temp_config"
}

# Function to activate profile
activate_profile() {
local profile_name=$1

if [ -z "$profile_name" ]; then
    echo "❌ Please specify a profile name"
    list_profiles
    return 1
fi

local config_file="$KUBE_DIR/config-$profile_name"

if [ ! -f "$config_file" ]; then
    echo "🔨 Profile config not found, building..."
    build_config "$profile_name"
fi

# Backup current config
cp "$KUBE_DIR/config" "$KUBE_DIR/config.previous" 2>/dev/null || true

# Activate profile
cp "$config_file" "$KUBE_DIR/config"

echo "✅ Activated profile: $profile_name"
echo "📊 Available contexts:"
kubectl config get-contexts

# Update shell prompt
export KUBECONFIG_PROFILE="$profile_name"
export PS1="[k8s:$profile_name] \u@\h:\w\$ "
}

# Function to list profiles
list_profiles() {
echo "📋 Available Profiles:"
if [ -d "$PROFILES_DIR" ]; then
    for profile in "$PROFILES_DIR"/*.yaml; do
```

```

if [ -f "$profile" ]; then
    local name=$(basename "$profile" .yaml)
    local desc=$(grep "DESCRIPTION=" "$profile" 2>/dev/null | cut -d '=' -f2- | tr -d '')
    echo " $name - $desc"
fi
done
else
    echo " No profiles found"
fi
}

# Function to create new profile
create_profile() {
    local profile_name=$1
    local description=$2

    if [ -z "$profile_name" ]; then
        echo "Usage: create_profile <name> <description>"
        return 1
    fi

    cat << PROFILE > "$PROFILES_DIR/$profile_name.yaml"
# Profile: $profile_name
DESCRIPTION="$description"
CLUSTERS=""
USERS=""
CONTEXTS=""
DEFAULT_CONTEXT=""

# Add cluster/user/context references here
# Example:
# CLUSTERS="dev-cluster prod-cluster"
# USERS="dev-user prod-user"
# CONTEXTS="dev-context prod-context"
# DEFAULT_CONTEXT="dev-context"
PROFILE

    echo " ✅ Created profile template: $profile_name"
    echo " 📄 Edit: $PROFILES_DIR/$profile_name.yaml"
}

# Function to validate configuration
validate_config() {
    local config_file=${1:-"$KUBE_DIR/config"}

```

```
echo "🔍 Validating kubeconfig: $config_file"

# Check if file exists
if [ ! -f "$config_file" ]; then
    echo "✗ Config file not found: $config_file"
    return 1
fi

# Validate YAML syntax
if ! kubectl config view --kubeconfig="$config_file" >/dev/null 2>&1; then
    echo "✗ Invalid YAML syntax"
    return 1
fi

# Check for required fields
local temp_kubeconfig="$config_file"
export KUBECONFIG="$temp_kubeconfig"

echo "📊 Configuration Summary:"
echo "Clusters: $(kubectl config get-clusters --no-headers | wc -l)"
echo "Users: $(kubectl config get-users --no-headers | wc -l)"
echo "Contexts: $(kubectl config get-contexts --no-headers | wc -l)"
echo "Current Context: $(kubectl config current-context 2>/dev/null || echo 'none')"

# Test connectivity to current context
if kubectl cluster-info --request-timeout=5s >/dev/null 2>&1; then
    echo "✓ Cluster connectivity: OK"
else
    echo "⚠️ Cluster connectivity: Failed (may be expected for mock configs)"
fi

echo "✓ Configuration validation complete"
}

# Function to show detailed context info
show_context_info() {
    local context_name=${1:-$(kubectl config current-context 2>/dev/null)}

    if [ -z "$context_name" ]; then
        echo "✗ No context specified or current context found"
        return 1
    fi
```

```

echo "🔍 Context Information: $context_name"

# Get context details
local cluster=$(kubectl config view -o jsonpath=".contexts[?(@.name=='$context_name)].context.cluster")
local user=$(kubectl config view -o jsonpath=".contexts[?(@.name=='$context_name)].context.user")
local namespace=$(kubectl config view -o jsonpath=".contexts[?(@.name=='$context_name)].context.namespace")

echo " Cluster: $cluster"
echo " User: $user"
echo " Namespace: ${namespace:-default}"

# Get cluster endpoint
local endpoint=$(kubectl config view -o jsonpath=".clusters[?(@.name=='$cluster')].cluster.server")
echo " Endpoint: $endpoint"

# Test permissions
echo ""
echo "🔒 Permission Check:"
kubectl auth can-i get pods --context="$context_name" 2>/dev/null && echo " ✅ Can get pods" || echo " ❌ Can't get pods"
kubectl auth can-i create deployments --context="$context_name" 2>/dev/null && echo " ✅ Can create deployments" || echo " ❌ Can't create deployments"
kubectl auth can-i '*' '*' --context="$context_name" 2>/dev/null && echo " ✅ Has admin access" || echo " ⚠️ Limited access"
}

# Main command dispatcher
case "$1" in
    "backup")
        create_backup
        ;;
    "build")
        build_config $2
        ;;
    "activate")
        activate_profile $2
        ;;
    "list")
        list_profiles
        ;;
    "create")
        create_profile $2 "$3"
        ;;
    "validate")
        validate_config $2
        ;;
    "info")
        ;;
esac

```

```

show_context_info $2
;;
"clean")
    echo " ✎ Cleaning up temporary files..."
    rm -f /tmp/kubeconfig-build-*
    echo " ✅ Cleanup complete"
;;
*)
    echo " 🛡 Advanced Kubeconfig Manager"
    echo """
    echo "Commands:"
    echo " backup           - Create full backup"
    echo " build <profile>     - Build config from profile"
    echo " activate <profile>   - Activate profile"
    echo " list              - List available profiles"
    echo " create <name> <desc>      - Create new profile"
    echo " validate [config]     - Validate configuration"
    echo " info [context]       - Show context information"
    echo " clean             - Clean temporary files"
    echo """
    echo "Examples:"
    echo " $0 create dev-profile 'Development environment'"
    echo " $0 build dev-profile"
    echo " $0 activate dev-profile"
;;
esac
EOF

chmod +x ~/kube/advanced-kubeconfig-manager.sh

# Create sample user profiles
create_user_profile "dev-user" "dev@company.com" "developer"
create_user_profile "ops-user" "ops@company.com" "admin"
create_user_profile "readonly-user" "viewer@company.com" "viewer"

# Create sample cluster profiles
create_cluster_profile "dev-cluster" "development" "https://dev-k8s.company.com" "us-west-2"
create_cluster_profile "staging-cluster" "staging" "https://staging-k8s.company.com" "us-east-1"
create_cluster_profile "prod-cluster" "production" "https://prod-k8s.company.com" "us-east-1"

# Create sample context templates
create_context_template "dev-context" "dev-cluster" "dev-user" "development" "Development environment context"
create_context_template "staging-context" "staging-cluster" "ops-user" "staging" "Staging environment context"
create_context_template "prod-context" "prod-cluster" "ops-user" "production" "Production environment context"

```

```
# Create sample profiles
cat << 'EOF' > ~/kube/profiles/developer.yaml
# Profile: developer
DESCRIPTION="Developer profile with access to dev and staging"
CLUSTERS="dev-cluster staging-cluster"
USERS="dev-user"
CONTEXTS="dev-context staging-context"
DEFAULT_CONTEXT="dev-context"
EOF

cat << 'EOF' > ~/kube/profiles/operator.yaml
# Profile: operator
DESCRIPTION="Operations profile with full access"
CLUSTERS="dev-cluster staging-cluster prod-cluster"
USERS="ops-user"
CONTEXTS="dev-context staging-context prod-context"
DEFAULT_CONTEXT="staging-context"
EOF

cat << 'EOF' > ~/kube/profiles/readonly.yaml
# Profile: readonly
DESCRIPTION="Read-only access to all environments"
CLUSTERS="dev-cluster staging-cluster prod-cluster"
USERS="readonly-user"
CONTEXTS="dev-context staging-context prod-context"
DEFAULT_CONTEXT="dev-context"
EOF

# Create convenient wrapper script
cat << 'EOF' > ~/kube/kconfig
#!/bin/bash
# Convenient wrapper for advanced kubeconfig manager
~/kube/advanced-kubeconfig-manager.sh "$@"
EOF

chmod +x ~/kube/kconfig

# Add to PATH
echo 'export PATH="$HOME/.kube:$PATH"' >> ~/.bashrc

echo "✅ Advanced Kubeconfig Management System Setup Complete!"
echo ""
echo "📁 Directory Structure:"
```

```
find ~/kube -type f -name "*.yaml" -o -name "*.sh" -o -name "kconfig" | head -15
```

```
echo ""  
echo "🚀 Usage Examples:"  
echo " kconfig list          # List available profiles"  
echo " kconfig create myprofile 'desc' # Create new profile"  
echo " kconfig build developer    # Build developer profile"  
echo " kconfig activate developer # Activate developer profile"  
echo " kconfig info            # Show current context info"  
echo " kconfig backup          # Create full backup"  
  
echo ""  
echo "📝 Next Steps:"  
echo "1. Edit profile files in ~/kube/profiles/"  
echo "2. Build and activate a profile"  
echo "3. Verify with: kubectl config get-contexts"
```

## 21. Advanced Tips

This final section covers advanced kubectl techniques, performance optimization, and production best practices.

### 21.1 Advanced kubectl Techniques

```
bash
```

```
# JSON Path queries - extract specific data from resources
kubectl get pods -o jsonpath='{.items[*].metadata.name}'
# Get all pod names

kubectl get pods -o jsonpath='{.items[*].status.podIP}'
# Get all pod IP addresses

kubectl get nodes -o jsonpath='{.items[*].status.addresses[?(@.type=="ExternalIP")].address}'
# Get external IP addresses of nodes

# Custom columns - create custom output formats
kubectl get pods -o custom-columns=NAME:.metadata.name,STATUS:.status.phase,NODE:.spec.nodeName,IP:.status.podIP

kubectl get deployments -o custom-columns=NAME:.metadata.name,REPLICAS:.spec.replicas,READY:.status.readyReplicas

# Field selectors - filter resources by field values
kubectl get pods --field-selector=status.phase=Running
kubectl get pods --field-selector=spec.nodeName=worker-1
kubectl get events --field-selector=type=Warning

# Label selectors - complex label-based filtering
kubectl get pods -l 'environment in (production,staging)'
kubectl get pods -l 'app=web,version!=v1'
kubectl get pods -l 'tier,environment notin (test)'

# Sort resources by various fields
kubectl get pods --sort-by=.metadata.creationTimestamp
kubectl get pods --sort-by=.status.startTime
kubectl get nodes --sort-by=.metadata.name

# Watch resources with specific output formats
kubectl get pods -w -o custom-columns=NAME:.metadata.name,STATUS:.status.phase

# Server-side apply - more reliable than client-side apply
kubectl apply --server-side -f deployment.yaml

# Dry-run with diff - see what would change before applying
kubectl diff -f deployment.yaml
kubectl apply --dry-run=server -f deployment.yaml

# Resource usage and limits
kubectl describe nodes | grep -A 5 "Allocated resources"
```

```
kubectl top pods --sort-by=cpu --all-namespaces  
kubectl top pods --sort-by=memory --containers
```

## Example: Advanced Resource Management and Automation

Create a comprehensive resource management and automation system:

```
bash
```

```

#!/bin/bash
# save as advanced-k8s-automation.sh
echo "🚀 Advanced Kubernetes Automation and Management"

# Create comprehensive resource analyzer
cat << 'EOF' > k8s-resource-analyzer.sh
#!/bin/bash

# Advanced resource analysis with multiple metrics
analyze_cluster_resources() {
    echo "📊 Comprehensive Cluster Resource Analysis"
    echo "====="

    # Node resource analysis
    echo "💻 Node Resources:"
    kubectl describe nodes | awk '
/^Name:/ { node=$2 }
/cpu:/ {
    if ($1 ~ /Allocated/) {
        allocated_cpu=$4
        total_cpu=$6
        cpu_percent=$7
        gsub(/[%]/, "", cpu_percent)
    }
}
/memory:/ {
    if ($1 ~ /Allocated/) {
        allocated_mem=$4
        total_mem=$6
        mem_percent=$7
        gsub(/[%]/, "", mem_percent)
        printf " %s: CPU %s/%s (%s%), Memory %s/%s (%s%)\n",
            node, allocated_cpu, total_cpu, cpu_percent,
            allocated_mem, total_mem, mem_percent
    }
}
    '
    echo ""

    echo "📦 Pod Resource Distribution:"
    kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.spec.containers[]?.resources.requests) |
"\(.metadata.namespace)\(\.metadata.name): CPU=\(.spec.containers[0].resources.requests.cpu // "none"), Memory="
'
}

```

```

' | head -10

echo ""
echo "⚠️ Pods Without Resource Limits:"
kubectl get pods --all-namespaces -o json | jq -r '.
.items[] |
select(.spec.containers[0].resources.limits == null) |
"\(.metadata.namespace)\(.metadata.name)"
' | head -5

echo ""
echo "🔥 High Resource Usage Pods:"
if kubectl top pods --all-namespaces >/dev/null 2>&1; then
    kubectl top pods --all-namespaces --sort-by=cpu | head -5
    echo ""
    kubectl top pods --all-namespaces --sort-by=memory | head -5
else
    echo " Metrics server not available"
fi
}

# Network analysis
analyze_network() {
    echo ""
    echo "🌐 Network Analysis:"
    echo "====="

    echo "📡 Services and Endpoints:"
    kubectl get services --all-namespaces -o wide | head -10

    echo ""
    echo "🔗 Ingress Resources:"
    kubectl get ingress --all-namespaces 2>/dev/null | head -5 || echo " No ingress resources found"

    echo ""
    echo "🔒 Network Policies:"
    kubectl get networkpolicies --all-namespaces 2>/dev/null | head -5 || echo " No network policies found"
}

# Storage analysis
analyze_storage() {
    echo ""
    echo "💾 Storage Analysis:"
    echo "====="

```

```

echo "● Persistent Volumes:"
kubectl get pv -o custom-columns=NAME:.metadata.name,CAPACITY:.spec.capacity,STATUS:.status.phase,CL

echo ""
echo "● Persistent Volume Claims:"
kubectl get pvc --all-namespaces -o custom-columns=NAMESPACE:.metadata.namespace,NAME:.metadata.name,ST

echo ""
echo "● Storage Classes:"
kubectl get storageclass -o custom-columns=NAME:.metadata.name,PROVISIONER:.provisioner,RECLAIMPO
}

# Security analysis
analyze_security() {
    echo ""
    echo "● Security Analysis:"
    echo "====="

    echo "● Service Accounts:"
    kubectl get serviceaccounts --all-namespaces | head -10

    echo ""
    echo "● RBAC Roles:"
    kubectl get roles,clusterroles --all-namespaces 2>/dev/null | head -10

    echo ""
    echo "● Secrets:"
    kubectl get secrets --all-namespaces -o custom-columns=NAMESPACE:.metadata.namespace,NAME:.metadata.name
}

# Performance analysis
analyze_performance() {
    echo ""
    echo "● Performance Analysis:"
    echo "====="

    echo "● Pod Restart Analysis:"
    kubectl get pods --all-namespaces -o json | jq -r '
.items[] |
select(.status.containerStatuses[]?.restartCount > 0) |
"\(.metadata.namespace)\(\.metadata.name): \(.status.containerStatuses[0].restartCount) restarts"
' | head -5
}

```

```

echo ""
echo "⌚ Long-running Pods:"
kubectl get pods --all-namespaces --sort-by=.metadata.creationTimestamp | head -5

echo ""
echo "⚠️ Failed Pods:"
kubectl get pods --all-namespaces --field-selector=status.phase=Failed 2>/dev/null | head -5
}

# Main execution
case "$1" in
"resources")
    analyze_cluster_resources
;;
"network")
    analyze_network
;;
"storage")
    analyze_storage
;;
"security")
    analyze_security
;;
"performance")
    analyze_performance
;;
*)
    analyze_cluster_resources
    analyze_network
    analyze_storage
    analyze_security
    analyze_performance
;;
esac
EOF

```

[chmod +x k8s-resource-analyzer.sh](#)

```

# Create cluster health monitor
cat << 'EOF' > k8s-health-monitor.sh
#!/bin/bash

```

```

ALERT_THRESHOLD_CPU=80
ALERT_THRESHOLD_MEMORY=80
ALERT_THRESHOLD_DISK=85
LOG_FILE="/tmp/k8s-health-monitor.log"

log_message() {
    echo "$(date '+%Y-%m-%d %H:%M:%S') - $1" | tee -a "$LOG_FILE"
}

check_node_health() {
    log_message "🖥️ Starting Node Health Check"

    while read -r node status; do
        if [[ "$status" != "Ready" ]]; then
            log_message "🔴 ALERT: Node $node is not ready (Status: $status)"
        else
            log_message "✅ Node $node is healthy"
        fi
    done <<(kubectl get nodes --no-headers -o custom-columns=NAME:.metadata.name,STATUS:.status.conditions[-1])
}

check_pod_health() {
    log_message "🔍 Starting Pod Health Check"

    # Check for pods in problematic states
    local problematic_pods=$(kubectl get pods --all-namespaces --no-headers | grep -E "(Error|CrashLoopBackOff|ImagePullBackOff|Pending)" | head -5 | while
        log_message " $line"
        done
    else
        log_message "✅ All pods are healthy"
    fi
}

check_resource_usage() {
    log_message "📊 Starting Resource Usage Check"

    if kubectl top nodes >/dev/null 2>&1; then
        # Check node resource usage
        kubectl top nodes --no-headers | while read node cpu_usage cpu_percent memory_usage memory_percent; do

```

```

cpu_num=$(echo $cpu_percent | tr -d '%')
mem_num=$(echo $memory_percent | tr -d '%')

if [ "$cpu_num" -gt "$ALERT_THRESHOLD_CPU" ]; then
    log_message "⚠️ ALERT: High CPU usage on $node: $cpu_percent"
fi

if [ "$mem_num" -gt "$ALERT_THRESHOLD_MEMORY" ]; then
    log_message "⚠️ ALERT: High memory usage on $node: $memory_percent"
fi

done
else
    log_message "⚠️ Metrics server not available - cannot check resource usage"
fi
}

check_storage_health() {
    log_message "💾 Starting Storage Health Check"

    # Check PVC status
    local pending_pvcs=$(kubectl get pvc --all-namespaces --no-headers | grep -c "Pending" || echo "0")
    if [ "$pending_pvcs" -gt 0 ]; then
        log_message "⚠️ ALERT: $pending_pvcs PVCs in Pending state"
    fi

    # Check PV availability
    local available_pvs=$(kubectl get pv --no-headers | grep -c "Available" || echo "0")
    log_message "ℹ️ Available PVs: $available_pvs"
}

check_service_health() {
    log_message "🌐 Starting Service Health Check"

    # Check services with no endpoints
    kubectl get endpoints --all-namespaces -o json | jq -r '
.items[] |
select(.subsets == null or .subsets == [])
"\(.metadata.namespace)\(.metadata.name)"
' | while read service; do
    if [ ! -z "$service" ]; then
        log_message "⚠️ ALERT: Service $service has no endpoints"
    fi
done
}

```

```
generate_health_report() {
    local report_file="/tmp/k8s-health-report-$(date +%Y%m%d-%H%M%S).html"

    cat << 'HTML' > "$report_file"
<!DOCTYPE html>
<html>
<head>
    <title>Kubernetes Health Report</title>
    <style>
        body { font-family: Arial, sans-serif; margin: 20px; }
        .healthy { color: green; }
        .warning { color: orange; }
        .error { color: red; }
        .section { margin: 20px 0; padding: 15px; border: 1px solid #ddd; }
        pre { background: #f5f5f5; padding: 10px; overflow-x: auto; }
    </style>
</head>
<body>
    <h1> 📊 Kubernetes Cluster Health Report</h1>
    <p>Generated: $(date)</p>

    <div class="section">
        <h2> 📊 Cluster Overview</h2>
        <pre>$(kubectl get nodes -o wide)</pre>
    </div>

    <div class="section">
        <h2> 🔍 Pod Status Summary</h2>
        <pre>$(kubectl get pods --all-namespaces | head -20)</pre>
    </div>

    <div class="section">
        <h2> ⚠️ Recent Alerts</h2>
        <pre>$(tail -20 "$LOG_FILE" 2>/dev/null || echo "No alerts logged")</pre>
    </div>

    <div class="section">
        <h2> 📈 Resource Usage</h2>
        <pre>$(kubectl top nodes 2>/dev/null || echo "Metrics not available")</pre>
    </div>
</body>
</html>
HTML
```

```
log_message "📋 Health report generated: $report_file"
echo "Health report: $report_file"
}

# Main monitoring loop
monitor_cluster() {
    log_message "🚀 Starting Kubernetes Health Monitor"

    while true; do
        check_node_health
        check_pod_health
        check_resource_usage
        check_storage_health
        check_service_health

        log_message "😴 Sleeping for 60 seconds..."
        sleep 60
    done
}

# Command dispatcher
case "$1" in
    "nodes")
        check_node_health
        ;;
    "pods")
        check_pod_health
        ;;
    "resources")
        check_resource_usage
        ;;
    "storage")
        check_storage_health
        ;;
    "services")
        check_service_health
        ;;
    "report")
        generate_health_report
        ;;
    "monitor")
        monitor_cluster
        ;;
    *)
```

```
*)  
echo "kube Kubernetes Health Monitor"  
echo "Usage: $0 {nodes|pods|resources|storage|services|report|monitor}"  
echo ""  
echo "Commands:"  
echo " nodes - Check node health"  
echo " pods - Check pod health"  
echo " resources - Check resource usage"  
echo " storage - Check storage health"  
echo " services - Check service health"  
echo " report - Generate HTML health report"  
echo " monitor - Start continuous monitoring"  
;;  
esac  
EOF
```

[chmod +x k8s-health-monitor.sh](#)

```
echo "✅ Advanced resource analyzer created!"  
echo "✅ Health monitor created!"
```

```
# Run the analyzers  
echo ""  
echo "🔍 Running cluster analysis..."  
../k8s-resource-analyzer.sh resources
```

```
echo ""  
echo "kube Running health check..."  
../k8s-health-monitor.sh pods
```

## Demo: Production-Ready Deployment Pipeline

Create a complete production deployment pipeline with all best practices:

```
bash
```

```
#!/bin/bash
# save as production-deployment-pipeline.sh
echo "─ ─ Production-Ready Kubernetes Deployment Pipeline"

# Create production-grade deployment template
cat << 'EOF' > production-app-template.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: ${APP_NAME}-${ENVIRONMENT}
  labels:
    environment: ${ENVIRONMENT}
    managed-by: deployment-pipeline
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: ${APP_NAME}-sa
  namespace: ${APP_NAME}-${ENVIRONMENT}
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: ${APP_NAME}-config
  namespace: ${APP_NAME}-${ENVIRONMENT}
data:
  app.properties: |
    environment=${ENVIRONMENT}
    log.level=${LOG_LEVEL}
    database.host=${DB_HOST}
    database.port=${DB_PORT}
    cache.enabled=${CACHE_ENABLED}
---
apiVersion: v1
kind: Secret
metadata:
  name: ${APP_NAME}-secrets
  namespace: ${APP_NAME}-${ENVIRONMENT}
type: Opaque
data:
  database-password: ${DB_PASSWORD_B64}
  api-key: ${API_KEY_B64}
---
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ${APP_NAME}
  namespace: ${APP_NAME}-${ENVIRONMENT}
  labels:
    app: ${APP_NAME}
    version: ${APP_VERSION}
    environment: ${ENVIRONMENT}
spec:
  replicas: ${REPLICAS}
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 25%
      maxSurge: 25%
  selector:
    matchLabels:
      app: ${APP_NAME}
  template:
    metadata:
      labels:
        app: ${APP_NAME}
        version: ${APP_VERSION}
        environment: ${ENVIRONMENT}
    annotations:
      prometheus.io/scrape: "true"
      prometheus.io/port: "8080"
      prometheus.io/path: "/metrics"
  spec:
    serviceAccountName: ${APP_NAME}-sa
    securityContext:
      runAsNonRoot: true
      runAsUser: 1000
      fsGroup: 2000
    containers:
      - name: ${APP_NAME}
        image: ${IMAGE_REGISTRY}/${APP_NAME}:${APP_VERSION}
        imagePullPolicy: Always
      ports:
        - containerPort:
```